

# Fondamenti di Data Science e Machine Learning

## Dimensionality Reduction

Aurelien Geron: «Hands on Machine Learning with Scikit Learn and TensorFlow, O'Reilly ed.

Prof. Giuseppe Polese, aa 2024-25

# Outline

- ▶ The Curve of Dimensionality
- ▶ Main Approaches for Dimensionality Reduction
  - ▶ Projection
  - ▶ Manifold Learning
- ▶ A brief introduction to PCA
  - ▶ Preserving the Variance
  - ▶ Principal Components
  - ▶ Projecting Down to  $d$  Dimensions
  - ▶ Using Scikit-Learn
  - ▶ Explained Variance Ratio
  - ▶ Choosing the Right Number of Dimensions

# Introduction


- ▶ Many Machine Learning problems involve thousands or even millions of features for each training instance
  - ▶ this make training extremely slow
  - ▶ it can also make it much harder to find a good solution
- ▶ This problem is often referred as **curse of dimensionality**
- ▶ In real-world problems
  - ▶ it is often possible to reduce the number of features considerably
    - ▶ Turning an intractable problem into a tractable one

# Reducing Dimensionality

- ▶ Reducing dimensionality does lose some information
  - ▶ just like compressing an image to JPEG can degrade its quality
- ▶ Even if it will speed up training, it may also make our system perform slightly worse
- ▶ It also makes our pipelines a bit more complex and thus harder to maintain
  - ▶ we should first try to train our system on the original data before considering dimensionality reduction if training is slow
- ▶ Reducing the dimensionality of training data may filter out some noise and thus result in higher performances
  - ▶ but in general it won't; it will just speed up training

# The Curve of Dimensionality (1)

- ▶ In general, our intuition fails us when we try to imagine a dimensional space higher than 3 dimensions
  - ▶ Even a basic 4D hypercube is incredibly hard to picture in our mind
  - ▶ Let alone a 200-dimensional ellipsoid bent in a 1,000-dimensional space



*Point, segment, square, cube, and tesseract (0D to 4D hypercubes)*

# The Curve of Dimensionality (2)

- ▶ It turns out that many things behave very differently in high-dimensional space
  - ▶ If you pick a random point in a unit square (a  $1 \times 1$  square), it will have only about a 0.4% chance of being located less than 0.001 from a border
    - ▶ it is very unlikely that a random point will be “extreme” along any dimension
- ▶ But in a 10,000-dimensional unit hypercube (a  $1 \times 1 \times \dots \times 1$  cube, with ten thousand 1s), this probability is greater than 99.999999%
- ▶ Most points in a high-dimensional hypercube are very close to the border

# The Curve of Dimensionality

- ▶ A more troublesome difference: if we randomly pick two points in a unit square, the distance between them will be, on average, roughly 0.52
- ▶ The average distance between two random points randomly picked in a unit 3D cube will be roughly 0.66.
- ▶ What about 2 points randomly picked in a 1,000,000-dimensional hypercube? Their average distance will be about 408.25 (rough  $\sqrt{1,000,000/6}$ ! )
- ▶ This is counterintuitive: how can 2 points be so far apart when both lie within the same unit hypercube?

# The Curve of Dimensionality (3)

- ▶ High-dimensional datasets risk of being very sparse
  - ▶ most training instances are likely to be far away from each other
- ▶ This also means that a new instance will likely be far away from any training instance, making predictions much less reliable than in lower dimensions, since they will be based on much larger extrapolations
  - ▶ In short, the more dimensions the training set has, the greater the risk of overfitting it



# The Curve of Dimensionality (4)


- ▶ In theory, one solution to the curse of dimensionality could be to increase the size of the training set to reach a sufficient density of training instances
- ▶ In practice, the number of training instances required to reach a given density grows exponentially with the number of dimensions
- ▶ Two main approaches to reducing dimensionality are
  - ▶ Projection
  - ▶ ManifoldLearning

# Projection (1)

- ▶ In most real-world problems training instances are not spread out uniformly across all dimensions
  - ▶ Many features are almost constant, while others are highly correlated
- ▶ As a result, all training instances actually lie within (or close to) a much lower-dimensional *subspace* of the high-dimensional space
- ▶ Let's look at an example
  - ▶ A 3D dataset represented by the circles (next slide)


# Projection (2)

- ▶ Notice that all training instances lie close to a plane
  - ▶ This is a lower-dimensional (2D) subspace of the high-dimensional (3D) space




# Projection (3)

- ▶ If we project every training instance perpendicularly onto this subspace
  - ▶ represented by the short lines connecting the instances to the plane
- ▶ we get the new 2D data
  - ▶ Note that the axes correspond to new features  $z_1$  and  $z_2$  (the coordinates of the projections on the plane)



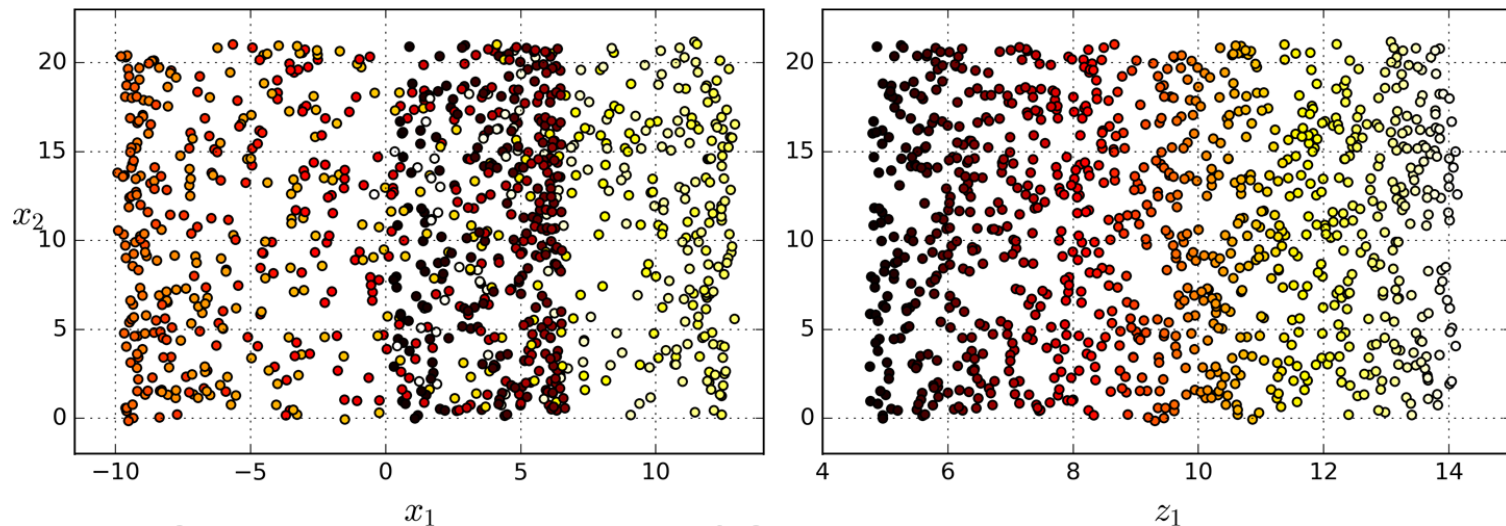
# Projection: Another Case (1)

- ▶ Projection is not always the best approach to dimensionality reduction
- ▶ In many cases the subspace may twist and turn, such as in the famous **Swiss roll** toy dataset



# Projection: Another Case (2)

- ▶ Simply projecting onto a plane (e.g., by dropping  $x_3$ ) would squash different layers of the Swiss roll together,
  - ▶ As shown on the left of Figure
- ▶ However, what you really want is to unroll the Swiss roll to obtain the 2D dataset
  - ▶ As shown on the right of Figure



# Manifold Learning (1)

- ▶ The Swiss roll is an example of a 2D manifold
  - ▶ Put simply, a 2D manifold is a 2D shape that can be bent and twisted in a higher-dimensional space
- ▶ More generally, a  $d$ -dimensional manifold is a part of an  $n$ -dimensional space (where  $d < n$ ) that locally resembles a  $d$ -dimensional hyperplane
- ▶ In the case of the Swiss roll,  $d = 2$  and  $n = 3$ 
  - ▶ It locally resembles a 2D plane, but it is rolled in the third dimension

# Manifold Learning (2)

- ▶ Many dimensionality reduction algorithms work by modeling the manifold on which the training instances lie
  - ▶ This is called **Manifold Learning**
- ▶ It relies on the manifold assumption, also called the manifold hypothesis
  - ▶ most real-world high-dimensional datasets lie close to a much lower-dimensional manifold
- ▶ This assumption is very often empirically observed



# Manifold Learning (3)

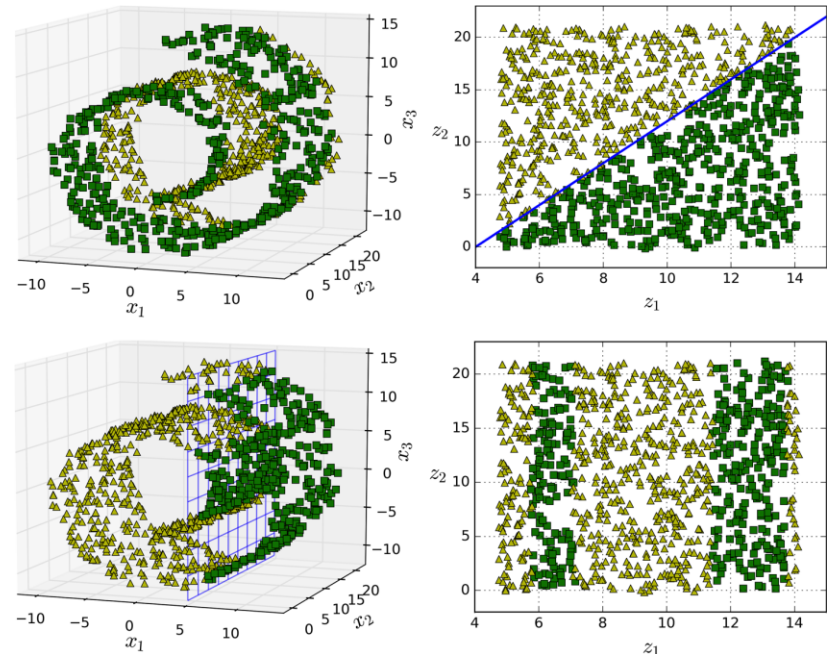
- ▶ Think about the MNIST dataset
  - ▶ All handwritten digit images have some similarities
- ▶ They are made of connected lines, the borders are white, they are more or less centered, and so on
- ▶ If we randomly generated images, only a tiny fraction of them would look like handwritten digits
  - ▶ The degrees of freedom available to us if we try to create a digit image are dramatically lower than the degrees of freedom we would have if we were allowed to generate any image we wanted
- ▶ These constraints tend to squeeze the dataset into a lower-dimensional manifold

# Manifold Learning (4)

- ▶ Another implicit assumption
  - ▶ The task at hand (e.g., classification or regression) will be simpler if expressed in the lower-dimensional space of the manifold
- ▶ In the top row of Figure (next slide) the Swiss roll is split into two classes
  - ▶ In the 3D space (on the left), the decision boundary would be fairly complex,
  - ▶ In the 2D unrolled manifold space (on the right), the decision boundary is a simple straight line

# Manifold Learning (5)

- ▶ The previous assumption does not always hold
  - ▶ In the bottom row of Figure the decision boundary is located at  $x_1 = 5$
- ▶ This decision boundary looks very simple in the original 3D space (a vertical plane), but it looks more complex in the unrolled manifold (a collection of four independent line segments)




# Summarizing Manifold Learning

- ▶ If we reduce the dimensionality of our training set before training a model
  - ▶ It will definitely speed up training
  - ▶ It may not always lead to a better or simpler solution
  - ▶ It all depends on the dataset
- ▶ We have a good sense of what the curse of dimensionality is and how dimensionality reduction algorithms can fight it, especially when the manifold assumption holds


# PCA

- ▶ **Principal Component Analysis** (PCA) is by far the most popular dimensionality reduction algorithm
  - ▶ First it identifies the hyperplane that lies closest to the data
  - ▶ Then it projects the data onto it




# Preserving the Variance (1)

- ▶ Before projecting the training set onto a lower-dimensional hyperplane, you first need to choose the right hyperplane
- ▶ For example
  - ▶ A simple 2D dataset is represented on the left of Figure, along with three different axes (i.e., one-dimensional hyperplanes)
  - ▶ The result of the projection onto each of these axes (on the right)




# Preserving the Variance (2)



- ▶ As you can see from the figure on the right
  - ▶ the projection onto the solid line preserves the maximum variance
  - ▶ the projection onto the dotted line preserves very little variance,
  - ▶ the projection onto the dashed line preserves an intermediate amount of variance

# Preserving the Variance (3)




- ▶ We will first project on the axis preserving the max variance
- ▶ This is the one minimizing the mean square distance between the original dataset and its projection on the selected axis
- ▶ Then, we project on the axis preserving the max residual variance, which is not the dashed line, since some of its variance was in  $c_1$
- ▶ **This is the rather simple idea behind PCA**



# Principal Components (1)


- ▶ PCA identifies the axis that accounts for the largest amount of variance in the training set (the solid line  $\mathbf{c}_1$ )
- ▶ It also finds a second axis, orthogonal to the first one, that accounts for the largest amount of remaining variance (the dotted line  $\mathbf{c}_2$ )



- ▶ PCA would also find a third axis, orthogonal to both previous axes, and a fourth, a fifth, and so on—as many axes as the number of dimensions in the dataset

# Principal Components (2)

- ▶ The unit vector that defines the  $i^{\text{th}}$  axis is called the  $i^{\text{th}}$  **principal component** (PC)
- ▶ In the previous Figure, the 1<sup>st</sup> PC is  $\mathbf{c}_1$  and the 2<sup>nd</sup> PC is  $\mathbf{c}_2$
- ▶ In the following Figure the first two PCs are represented by the orthogonal arrows in the plane, and the third PC would be orthogonal to the plane (pointing up or down)



# Principal Components (3)

- ▶ The direction of the principal components is not stable
  - ▶ if you perturb the training set slightly and run PCA again, some of the new PCs may point in the opposite direction of the original PCs
- ▶ However, they will generally still lie on the same axes
- ▶ In some cases, a pair of PCs may even rotate or swap, but the plane they define will generally remain the same

# Principal Component Matrix

*How can you find the principal components of a training set?*

- ▶ There is a standard matrix factorization technique called **Singular Value Decomposition** (SVD)
- ▶ It can decompose the training set matrix  $\mathbf{X}$  into the dot product of three matrices  $\mathbf{U} \cdot \Sigma \cdot \mathbf{V}^T$ , where  $\mathbf{V}$  contains all the principal components that we are looking for

$$\mathbf{V} = \begin{pmatrix} | & | & \dots & | \\ c_1 & c_2 & \dots & c_n \\ | & | & & | \end{pmatrix}$$

# Projecting Down to $d$ Dimensions

- ▶ Once you have identified all the principal components
  - ▶ you can reduce the dimensionality of the dataset down to  $d$  dimensions by projecting it onto the hyperplane defined by the first  $d$  principal components
  - ▶ Selecting this hyperplane ensures that the projection will preserve as much variance as possible
  - ▶ You can simply compute the dot product of the training set matrix  $\mathbf{X}$  by the matrix  $\mathbf{W}_d$ 
    - ▶ defined as the matrix containing the first  $d$  principal components

$$\mathbf{X}_{d\text{-proj}} = \mathbf{X} \cdot \mathbf{W}_d$$

# Using Scikit-Learn for PCA (1)

- ▶ Scikit-Learn's PCA class implements PCA using SVD decomposition just like we did before
- ▶ The following code applies PCA to reduce the dimensionality of the dataset down to two dimensions

```
from sklearn.decomposition import PCA
import numpy as np
#Build 3D dataset:
np.random.seed(4)
m = 60
w1, w2 = 0.1, 0.3
noise = 0.1
angles = np.random.rand(m) * 3 * np.pi / 2 - 0.5
X = np.empty((m, 3))
X[:, 0] = np.cos(angles) + np.sin(angles)/2 + noise * np.random.randn(m) / 2
X[:, 1] = np.sin(angles) * 0.7 + noise * np.random.randn(m) / 2
X[:, 2] = X[:, 0] * w1 + X[:, 1] * w2 + noise * np.random.randn(m)
#PCA using Scikit-Learn
pca = PCA(n_components = 2)
X2D = pca.fit_transform(X)
print("X2D[:5]:\n{}".format(X2D[:5]))
```

```
X2D[:5]:
[[ 1.26203346  0.42067648]
 [-0.08001485 -0.35272239]
 [ 1.17545763  0.36085729]
 [ 0.89305601 -0.30862856]
 [ 0.73016287 -0.25404049]]
```

# Using Scikit-Learn for PCA (2)

- ▶ After fitting the PCA transformer to the dataset, you can access the principal components using the `components_` variable
- ▶ Note that it contains the PCs as horizontal vectors
  - ▶ for example, the first principal component is equal to `pca.components_.T[:, 0]`

# Explained Variance Ratio

- ▶ Another very useful piece of information is the **explained variance ratio** of each principal component
  - ▶ available via the `explained_variance_ratio_` variable
- ▶ It indicates the proportion of the dataset's variance that lies along the axis of each principal component

- ▶ For example

```
print("explained_variance:\n{}".format  
      (pca.explained_variance_ratio_))
```

```
explained_variance:  
[0.84248607 0.14631839]
```

```
Process finished with exit code 0
```

- ▶ This tells you that 84.2% of the dataset's variance lies along the first axis, and 14.6% lies along the second axis
- ▶ This leaves less than 1.2% for the third axis, so it is reasonable to assume that it probably carries little information



# The Right Number of Dimensions (1)

- ▶ It is generally preferable to choose the number of dimensions that add up to a sufficiently large portion of the variance (e.g., 95%)
  - ▶ Unless you are reducing dimensionality for data visualization
    - ▶ In that case you will generally want to reduce the dimensionality down to 2 or 3
- ▶ The following code computes PCA without reducing dimensionality, then computes the minimum number of dimensions required to preserve 95% of the training set's variance

# The Right Number of Dimensions (2)

```
from sklearn.decomposition import PCA
import numpy as np
from six.moves import urllib
#Use of MNIST dataset
try:
    from sklearn.datasets import fetch_openml
    mnist = fetch_openml('mnist_784', version=1)
    mnist.target = mnist.target.astype(np.int64)
except ImportError:
    from sklearn.datasets import fetch_mldata
    mnist = fetch_mldata('MNIST original')
from sklearn.model_selection import train_test_split
X = mnist["data"]
y = mnist["target"]
X_train, X_test, y_train, y_test = train_test_split(X, y)
#PCA using Scikit-Learn
pca = PCA()
pca.fit(X_train)
cumsum = np.cumsum(pca.explained_variance_ratio_)
d = np.argmax(cumsum >= 0.95) + 1
print("d: {}".format(d))
```

d:154

Process finished with exit code 0

# The Right Number of Dimensions (3)

- ▶ You could then set `n_components=d` and run PCA again
  - ▶ Instead of specifying the number of principal, you can set `n_components` to be a float between 0.0 and 1.0, indicating the ratio of variance you wish to preserve

```
from sklearn.decomposition import PCA
import numpy as np
from six.moves import urllib
#Use of MNIST dataset
try:
    from sklearn.datasets import fetch_openml
    mnist = fetch_openml('mnist_784', version=1)
    mnist.target = mnist.target.astype(np.int64)
except ImportError:
    from sklearn.datasets import fetch_mldata
    mnist = fetch_mldata('MNIST original')
from sklearn.model_selection import train_test_split
X = mnist["data"]
y = mnist["target"]
X_train, X_test, y_train, y_test = train_test_split(X, y)
#PCA using Scikit-Learn
pca = PCA(n_components=0.95)
X_reduced = pca.fit_transform(X_train)
print("pca.n_components_:{0}".format(pca.n_components_))
```

```
pca.n_components_:154
```

```
Process finished with exit code 0
```

# The Right Number of Dimensions (4)

- ▶ Yet another option is to plot the explained variance as a function of the number of dimensions
  - ▶ simply plot `cumsum`
  - ▶ There will usually be an elbow in the curve, where the explained variance stops growing fast
- ▶ You can think of this as the intrinsic dimensionality of the dataset
  - ▶ In this case, you can see that reducing the dimensionality down to about 100 dimensions wouldn't lose too much explained variance

