Quantum Cryptography: A Survey

DAGMAR BRUSS, GÁBOR ERDÉLYI, TIM MEYER, TOBIAS RIEGE, and JÖRG ROTHE

Heinrich-Heine-Universität Düsseldorf

We survey some results in quantum cryptography. After a brief introduction to classical cryptography, we provide the quantum-mechanical background needed to present some fundamental protocols from quantum cryptography. In particular, we review quantum key distribution via the BB84 protocol and its security proof, as well as the related quantum bit commitment protocol and its proof of insecurity.

1. INTRODUCTION

Cryptography is the science of keeping private information from unauthorized access, of ensuring data integrity and authentication, and other tasks. In this survey, we will focus on quantum-cryptographic key distribution and bit commitment protocols and we in particular will discuss their security. Before turning to quantum cryptography, let us give a brief review of classical cryptography, its current challenges and its historical development.

Two parties, Alice and Bob, wish to exchange messages via some insecure channel in a way that protects their messages from eavesdropping. An algorithm, which is called a *cipher* in this context, scrambles Alice's message via some rule such that restoring the original message is hard—if not impossible—without knowledge of the secret key. This "scrambled" message is called the ciphertext. On the other hand, Bob (who possesses the secret key) can easily decipher Alice's ciphertext and obtains her original plaintext. Figure 2 in the next section presents this basic cryptographic scenario.

Cryptographic technology in use today relies on the hardness of certain mathematical problems. Classical cryptography faces the following two problems. First, the security of many classical cryptosystems is based on the hardness of problems such as integer factoring or the discrete logarithm problem. But since these problems typically are not *provably* hard, the corresponding cryptosystems are potentially insecure. For example, the famous and widely used RSA public-key cryptosystem [Rivest et al. 1978] could easily be broken if large integers were easy to factor. The hardness of integer factoring, however, is not a proven fact but rather a hypothesis.[1] We mention in passing that computing the RSA secret key from the corresponding public key is polynomial-time equivalent to integer factoring [May 2004].

Second, the theory of quantum computation has yielded new methods to tackle these mathematical problems in a much more efficient way. Although there are still numerous challenges to overcome before a working quantum computer of sufficient power can be built, in theory many classical ciphers (in particular public-key cryptosystems such as RSA) might be broken by such a powerful machine. However, while quantum computation seems to be a severe challenge to classical cryptography in a possibly not so distant future, at the same time it offers new possibilities to build encryption methods that are safe even against attacks performed by means of a quantum computer. Quantum cryptography extends the power of classical cryptography by protecting the secrecy of messages using the physical laws of quantum mechanics.

Looking back in the history of cryptography, one of the first encryption methods was the scytale. The first recorded use of the scytale dates back to the fifth century B.C. when the Spartans used it to exchange battle information between generals without revealing it to the enemy. To encrypt a message, called the *plaintext*, a strip of leather or pergament was wrapped around a wooden cylinder, the scytale. The encrypted message, also called the *ciphertext*, was then written from left to right onto the leather, so that unravelling the strip would produce a meaningless alignment of seemingly random letters, see Figure 1 for the encryption of the plaintext "scytaleisatranspositioncipher" by "ssoicaspytihtrteaairlnoesnipc." The decryption of the ciphertext was achieved by using a scytale of the same diameter as the cylinder that was used for encryption.

The scytale is a so-called transposition cipher, since only the order of the letters within the message is changed. Another type of encryption is the substitution cipher.

---

[1]Notwithstanding the fact that we currently do not have any (classical) polynomial-time algorithm for factoring integers, Fellows and Koblitz [1992] provided evidence that the (decision version of the) integer factoring problem in fact is far from being hard (in the traditional model of worst-case complexity), by showing that it is unlikely to be NP-complete. Of course, even if this problem were NP-complete, it still might happen that P = NP (which itself is a famous open question) and so all NP-complete problems would have polynomial-time algorithms. And even if integer factoring were NP-complete and P ≠ NP, it still might happen that integers could be factored in polynomial time *on the average*. For the average-case complexity model, we refer to Levin's work [Levin 1986] and to the surveys by Goldreich [1997] and Wang [1997]. On a related note, Ajtai and Dwork [1997] proposed a cryptosystem whose security is based on a lattice problem shown to be equally hard in the worst case and in the average case, see Nguyen and Stern [1998] for the cryptanalysis of this cryptosystem.

Fig. 1. The Scytale.

Here, instead of swapping the positions of the letters, each plaintext letter is replaced by another letter according to some specific rule.

The method of encryption and decryption is called a *cryptosystem*, whereas the particular information used for encryption or decryption in an individual communication is called a *key*. In the case of the scytale, the diameter of the cylinder represents the secret key. Obviously, this ancient cryptosystem has a very low level of security. Once the method of encryption is known to the eavesdropper, he or she can simply try all possible diameters to reveal the original message. The fact that the cryptosystem is publicly known is not the reason for the insecurity of the communication, but rather the small number of possible keys that can be used for encryption. In the 19th century, Auguste Kerckhoffs stated the principle that the security of a cryptosystem must be based solely on the secrecy of the key itself. Therefore, when designing new ciphers, one should always treat the algorithm as if it were publicly known.

Over time, the amount of information that needed to be encrypted exploded, making it impossible to use simple and insecure procedures like the scytale. At first, mechanical devices were built to speed up the encryption and the (authorized) decryption process, and to increase the complexity of the keys used to scramble the messages. An infamous example of such a mechanical cryptosystem is the Enigma, which was used in World War II by the Germans to conceal their military communication. Not being aware of certain weaknesses of their encryption device (the most significant of which was that this substitution cipher allowed for known-plaintext attacks), the Germans considered the Enigma unbreakable. However, allied cryptanalysts in Bletchley Park near London often were able to decrypt the Germans' military messages during the war. One might argue that breaking the Enigma was one of the most crucial factors for the victory of the allied forces and for ending the war. After the war, it was the invention of the transistor that made the rise of the computer industry possible.

The huge speed-up in executing mathematical calculations resulted in the need to create much more secure cryptosystems, among them symmetric block ciphers such as the Data Encryption Standard (DES) and the Advanced Encryption Standard (AES) and public-key cryptosystems such as RSA and others, which are integrated in modern cryptographic applications currently in use. A nice and easy-to-read overview of the history of cryptography is given by Singh [1999]. With the currently emerging theory of quantum computation, we seem to be at the beginning of yet another era of cryptography.

This survey is organized as follows. Section 2 describes the fundamentals of classical cryptography including an easy example. Section 3 provides some background of quantum mechanics and introduces our notation. In Section 4, we present the BB84 quantum key distribution protocol and discuss its security. In particular, we describe an entanglement-based version of BB84, which is akin to Ekert's protocol [Ekert 1991] protocol (see also Bennett et al. [1992]), provide a proof of security for this protocol, and show that it is equivalent to the "prepare-and-measure" version of the BB84 protocol. Section 5 presents the BB84 quantum bit commitment protocol and shows that the security of unconditional quantum bit commitment is impossible. Finally, Section 6 gives a brief outlook and draws some conclusions.

2. CLASSICAL CRYPTOGRAPHY

Overviews of classical cryptography can be found in various text books (see, e.g., Rothe [2005] and Stinson [2005]). Here, we present just the basic definition of a cryptosystem and give one example of a classical encryption method, the one-time pad.

Definition 2.1. A *(deterministic, symmetric) cryptosystem* is a five-tuple $(\mathcal{P}, \mathcal{C}, \mathcal{K}, \mathcal{E}, \mathcal{D})$ satisfying the following conditions:

(1) $\mathcal{P}$ is a finite set of possible *plaintexts*.
(2) $\mathcal{C}$ is a finite set of possible *ciphertexts*.
(3) $\mathcal{K}$ is a finite set of possible *keys*.
(4) For each $k \in \mathcal{K}$, there are an *encryption rule* $e_k \in \mathcal{E}$ and a corresponding *decryption rule* $d_k \in \mathcal{D}$, where $e_k : \mathcal{P} \to \mathcal{C}$ and $d_k : \mathcal{C} \to \mathcal{P}$ are functions satisfying $d_k(e_k(x)) = x$ for each plaintext element $x \in \mathcal{P}$.

In the basic scenario in cryptography, we have two parties who wish to communicate over an insecure channel, such as a phone line or a computer network. Usually, these parties are referred to as Alice and Bob. Since the communication channel is insecure, an eavesdropper, called Eve, may intercept the messages that are sent over this channel. By agreeing on a secret key k via a secure communication method, Alice and Bob can make use of a cryptosystem to keep their information secret, even when sent over the insecure channel. This situation is illustrated in Figure 2.

The method of encryption works as follows. For her secret message m, Alice uses the key k and the encryption rule e_k to obtain the ciphertext $c = e_k(m)$. She sends Bob the ciphertext c over the insecure channel. Knowing the key k, Bob can easily decrypt the ciphertext by the decryption rule d_k:

$$d_k(c) = d_k(e_k(m)) = m.$$

Knowing the ciphertext c but missing the key k, there is no easy way for Eve to determine the original message m.

There exist many cryptosystems in modern cryptography to transmit secret messages. An early well-known system is the *one-time pad*, which is also known as the

A	B	C	D	E	$\cdots$	X	Y	Z		!	–	.
00	01	02	03	04	$\cdots$	23	24	25	26	27	28	29

Fig. 3. Letters and punctuation marks encoded by numbers from 0 to 29.

plaintext m	O	N	E	–	T	I	M	E		P	A	D
m encoded	14	13	04	28	19	08	12	04	26	15	00	03
key k	06	13	02	01	14	05	07	18	05	26	13	28
c encoded	20	26	06	29	03	13	19	22	01	11	13	01
ciphertext c	U		G	.	D	N	T	W	B	L	N	B

Fig. 4. Encryption and decryption example for the one-time pad.

Vernam cipher. The one-time pad is a substitution cipher. Despite its advantageous properties, which we will discuss later on, the one-time pad's drawback is the costly effort needed to transmit and store the secret keys.

Example 2.2 (One-Time Pad). For plaintext elements in $\mathcal{P}$, we use capital letters and some punctuation marks, which we encode as numbers ranging from 0 to 29, see Figure 3. As is the case with most cryptosystems, the ciphertext space equals the plaintext space. Furthermore, the key space $\mathcal{K}$ also equals $\mathcal{P}$, and we have $\mathcal{P} = \mathcal{C} = \mathcal{K} = \{0, 1, \ldots, 29\}$.

Next, we describe how Alice and Bob use the one-time pad to transmit their messages. A concrete example is shown in Figure 4. Suppose Alice and Bob share a joint secret key k of length $n = 12$, where each key symbol $k_i \in \{0, 1, \ldots, 29\}$ is chosen uniformly at random. Let $m = m_1 m_2 \ldots m_n$ be a given message of length n, which Alice wishes to encrypt. For each plaintext letter m_i, where $1 \leq i \leq n$, Alice adds the plaintext numbers to the key numbers. The result is taken modulo 30. For example, the last letter of the plaintext from Figure 4, "D," is encoded by "$m_{12} = 03$." The corresponding key is "$m_{12} = 28$," so we have $c_{12} = 3 + 28 = 31$. Since $31 \equiv 1 \bmod 30$, our plaintext letter "D" is encrypted as "B." Decryption works similarly by subtracting, character by character, the key letters from the corresponding ciphertext letters. So the encryption and decryption can be written as respectively $c_i = (m_i + k_i) \bmod 30$ and $m_i = (c_i - k_i) \bmod 30$, $1 \leq i \leq n$.

We will prove that the one-time pad achieves perfect secrecy. To define perfect secrecy, we need some elementary notions from probability theory.

Notation 2.3. Let **X** be a discrete random variable that can take on values from a finite set $\mathcal{X}$ according to a given probability distribution on $\mathcal{X}$. We denote by $\Pr[\mathbf{X} = x]$ the probability that **X** takes on the value $x \in \mathcal{X}$. If **X** is clear from the context, we just write $\Pr[x]$. For all $x \in \mathcal{X}$, $\Pr[x] \geq 0$. Additionally, $\sum_{x \in \mathcal{X}} \Pr[x] = 1$. For another random variable **Y** defined on the finite set $\mathcal{Y}$, we denote by $\Pr[x|y]$ the conditional probability that **X** takes on the value $x \in \mathcal{X}$ given that **Y** takes on the value $y \in \mathcal{Y}$.

Suppose that a probability distribution on the finite plaintext space $\mathcal{P}$ is given. Thus, the plaintext element defines a random variable, which we denote by **p**. Similarly, the key chosen by Alice and Bob for their communication defines a random variable on the key space, denoted by **k**. Both probability distributions, for **p** and **k**, induce a probability distribution on the ciphertext space $\mathcal{C}$, which gives another random variable **c** for the ciphertext element. We now define the notion of perfect secrecy that was introduced by Shannon [1949].

Definition 2.4.   A cryptosystem is said to achieve *perfect secrecy* if and only if for each $p \in \mathcal{P}$ and for each $c \in \mathcal{C}$,

$$\Pr[p|c] = \Pr[p].$$

That means that the event that some plaintext p was encrypted is independent of the ciphertext c being observed. In other words, knowing c yields no advantage when trying to retrieve the original plaintext p.

In his pathbreaking paper, Shannon [1949] showed that for any cryptosystem achieving perfect secrecy, the uncertainty about the key used for encryption (as measured by the entropy of the key space) is at least as large as the uncertainty about the message encrypted. We here state a characterization of when perfect secrecy can be achieved, which also is sometimes referred to as Shannon's Theorem (and the proof of which can be found in, for example, Rothe [2005] and Stinson [2005]): Suppose that $(\mathcal{P}, \mathcal{C}, \mathcal{K}, \mathcal{E}, \mathcal{D})$ is a cryptosystem with $\|\mathcal{K}\| = \|\mathcal{P}\| = \|\mathcal{C}\|$ and such that every plaintext element will be encrypted with a positive probability. Then, this cryptosystem achieves perfect secrecy if and only if

(1) the keys in $\mathcal{K}$ are uniformly distributed, and

(2) for each $p \in \mathcal{P}$ and for each $c \in \mathcal{C}$, there exists a unique key k such that $e_k(p) = c$.

Using this characterization, it is easy to see that the one-time pad satisfies the property of perfect secrecy. Since a new key element is created for each single plaintext element randomly under the uniform distribution, knowing the ciphertext is no advantage for an eavesdropper who seeks to recover the original message.

In addition to providing perfect secrecy, the one-time pad allows the choice of timing: Keys are transmitted whenever possible, and then encryption is done whenever needed. However, the one-time pad also has severe disadvantages that make it impractical to use. Recall that the key has to be as large as the message itself. Thus, the number of bits that need to be exchanged over a secure channel for obtaining a joint secret key increases with the amount of information that Alice and Bob wish to transmit secretly. In light of this fact, one might ask why they don't use the secure channel directly for their communication. Using the same key for encryption more than once is no alternative, as the one-time pad's perfect secrecy crucially depends on creating a new key for every single plaintext element.

The scytale and the one-time pad are two examples of a symmetric cryptosystem. That means that the same key is used for encryption and decryption (or, at least, that the decryption key can be easily determined from the encryption key). Thus, in order to use such a cryptosystem, Alice and Bob first (i.e., prior to executing the protocol) have to agree on a joint secret key. Since the encryption and decryption keys (essentially) are the same, it might seem that this secret-key agreement necessarily requires an expensive secure channel. It may not immediately be obvious how two parties can agree on a joint secret key via communicating over an insecure (and inexpensive) public channel: If one party simply chooses some key and sends it encrypted to the other party, then which key should be used to encrypt the other key in the first place? This dilemma is known as the key agreement problem, and it was long considered unsolvable. However, Diffie and Hellman [1976] found a quite simple but brilliant way to avoid this dilemma and to solve the key agreement problem, making use of the hardness of the discrete logarithm problem. Other secret-key agreement protocols were proposed by Rivest and Sherman (see Rabi and Sherman [1997], Hemaspaandra and Rothe [1999], and Hemaspaandra et al. [2005, 2006]), and others.

Diffie and Hellman's secret-key agreement protocol enables Alice and Bob to agree on a joint secret key by communicating over a public channel, and even though Eve intercepts each bit transmitted she is not able to determine the secret key, provided

that discrete logarithms are hard to compute. As mentioned earlier, whether or not the discrete logarithm problem indeed is hard is an open question, and it is also not known whether or not computing discrete logarithms is as hard as breaking Diffie–Hellman (see Maurer and Wolf [1999] for more details).

A major disadvantage of symmetric ciphers and the related issue of key distribution occurs when many parties in a large communication network need to share joint secret keys. In principle, if n parties participate, $n(n + 1)/2$ different secret keys would have to be generated. Public-key cryptosystems, also called asymmetric cryptosystems, circumvent this key distribution problem as follows: Instead of having one key for every pair of parties, only one pair of keys per party is needed to communicate securely. Diffie and Hellman [1976] proposed the principle idea of public-key cryptography, namely to use two distinct keys, a public key for encryption and a private key for decryption.[2]

The first public-key cryptosystem that appeared in the open literature[3] is the RSA system, named after its three inventors, Ron Rivest, Adi Shamir, and Leonard Adleman [Rivest et al. 1978]. Up to date, RSA is still used in numerous cryptographic applications. Public-key cryptosystems are based on so-called *(trapdoor) one-way functions*, functions that are easy to compute but hard to invert (unless one possesses a certain "trapdoor" information required for authorized decryption).

To communicate via a public-key cryptosystem, Alice creates two keys, k_{public} and $k_{private}$. Her encryption key k_{public} is public, but Alice keeps her private decryption key $k_{private}$ secret. Each time Bob wishes to communicate with Alice, he looks up her public key and uses it to encrypt his message. Since only Alice knows her private key, she alone can (efficiently) decrypt the ciphertext, that is, invert the encryption function, which is one-way.

Unfortunately, it is a central open question whether one-way functions exist. The notion of one-way function has been intensely studied in various contexts. In particular, "noninvertibility" as implicit in the definition of one-way-ness strongly depends on the computational model used, and so do concrete candidates of one-way functions. Berman [1977], Brassard et al. [1978], Brassard [1979], Ko [1985], and—perhaps most notably—Grollmann and Selman [1988] were among the first to study one-way functions in the traditional model of worst-case complexity. Such "complexity-theoretic" one-way functions and one-way permutations have been further investigated in, for example, Allender and Rubinstein [1988], Watanabe [1988], Hartmanis and Hemachandra [1991], Selman [1992], Hemaspaandra et al. [1997a, 1997b, 2005, 2006], Rabi and Sherman [1997], Hemaspaandra and Rothe [1999, 2000], Beygelzimer et al. [1999], Rothe and Hemaspaandra [2002], Fenner et al. [2003], Homan and Thakur [2003], and Homan [2004]. Along a different path, one-way functions were carefully studied in the more challenging average-case complexity model, which is central to cryptographic applications. To mention just one result along this line of research, Håstad et al. [1999] have shown how to construct pseudorandom number generators not only from any one-way permutation but even from any given one-way function. Finally, the study of one-way functions in quantum cryptography, not surprisingly, was initiated not long ago. To mention just one recent result here, Kawachi et al. [2005] provided a necessary

---

[2]Interestingly, Diffie and Hellman [1976] simultaneously solved the key agreement problem and proposed public-key cryptography, which makes the use of secret-key agreement obsolete. Note, however, that symmetric cryptosystems do have important advantages, such as being more efficient than most public-key cryptosystems, which makes them and the corresponding key agreement protocols still very useful in practice.

[3]In 1997, the British Government Communications Headquarters revealed that its researchers James Ellis, William Cocks, and Malcolm Williamson had independently and even earlier discovered the principle idea of public-key cryptography, the cryptosystem now called RSA, and the secret-key agreement protocol now called Diffie–Hellman, see, for example, the discussion in Singh [1999] and Rothe [2005].

and sufficient condition that can be used as a universal test for quantum one-way functions and that is akin to the next-bit test for pseudorandom number generators.

The key issue is to find one-way functions that are secure enough to use for public-key cryptography. The first one-way function designed for this purpose (i.e., the RSA encryption function) is based on the problem of factoring large integers. As mentioned in the introduction, no efficient classical algorithm for computing the prime factors of some given integer is known up to now (see Footnote 1). Other public-key cryptosystems—such as the ElGamal system [ElGamal 1985]—are based on the presumed hardness of computing discrete logarithms. One disadvantage of such systems is that they typically lack a proof of security. Another disadvantage is that the directory storing the public keys has to be protected against manipulation and unauthorized access. If eavesdropper Eve replaces Alice's public key with her own key, she can decrypt all messages sent to Alice.

Since Peter Shor proposed his celebrated polynomial-time algorithms for factoring integers and computing discrete logarithms with quantum computers [Shor 1997], all cryptosystems whose security is based on the hardness of solving these mathematical problems have become (at least theoretically) vulnerable. Although it will certainly take some time for the first practical quantum computers to emerge, it is advisable to look for alternative, new cryptosystems whose security is not based solely on the hardness of solving such mathematical problems with current computer technology. Quantum theory seems to be the perfect basis on which to build such a new cryptosystem that withstands even an attack by quantum computers.

3. FROM BITS TO QUBITS

The most important unit of information in computer science is the *bit*. There are two possible values that can be stored by a bit: the bit is either equal to "0" or equal to "1." These two different states can be represented in various ways, for example by a simple switch or by a capacitor: if not charged, the capacitor holds the value zero; if charged, it holds the value one.

In general, a quantum state $|\psi\rangle$ is an element of a finite-dimensional complex vector space (or Hilbert space) H. We denote the scalar product of two states $|\psi\rangle$ and $|\phi\rangle$ by $\langle\psi|\phi\rangle$, where $\langle\psi| = \overline{|\psi\rangle}^T$ is the conjugate transpose of $|\psi\rangle$.[4] It is convenient to deal with normalized states, so we require $\langle\psi|\psi\rangle = 1$ for all states $|\psi\rangle$ that have a physical meaning.

The quantum analog of the bit is called *qubit*, which is derived from *qu*antum *bit*. A qubit $|\psi\rangle$ is an element of a two-dimensional Hilbert space, in which we can introduce an orthonormal basis, consisting of the two states $|0\rangle$ and $|1\rangle$. Unlike its classical counterpart, the quantum state can be in any *coherent superposition* of the basis states:

$$|\psi\rangle = \alpha|0\rangle + \beta|1\rangle, \tag{1}$$

where α and β are, in general, complex coefficients. This is due to the fact that the quantum mechanical equation of motion, the Schrödinger equation, is linear: Any linear superposition of its solutions (the quantum states) is also a solution. Since we require quantum states to be normalized, we find that the coefficients in (1) have to fulfill $|\alpha|^2 + |\beta|^2 = 1$, where $|\cdot|$ denotes the absolute value.

There exist many possibilities to physically represent a qubit in practice, as every quantum system with at least two states can serve as a qubit. For example, the spin of

---

[4]Mathematically, $\langle\psi|$ is an element of the dual space H^*.

an atom or the polarization[5] of a light particle can represent the state of a qubit. Even a cat with its two basic states "dead" and "alive," introduced by Schrödinger [1935] to visualize fundamental concepts of quantum mechanics, might serve as a representation. The cat's problem—or fortune from the animal's point of view—when being used as a quantum system is its sheer size compared to that of an atom or light particle. There is no way to protect such a big quantum instance from interaction with its environment, which in turn will result in decoherence of the superposition of the cat. For the rest of the chapter, we will leave the cat alone and use light particles as our preferred qubits.

The physical meaning of (1) can most easily be understood when we measure the quantum state $|\psi\rangle$. In quantum mechanics, this is achieved by a *positive operator valued measurement* (POVM), which is a family of positive-definite, hermitian operators $\mathcal{E} = \{E_x\}_{x \in \mathcal{X}}$ acting on the Hilbert space of the qubit. The members of this family have to sum up to the identity, $\sum_{x \in \mathcal{X}} E_x = \mathbb{1}$. A simple, special case occurs when the E_x are orthogonal projectors, that is, $E_x = |\phi_x\rangle\langle\phi_x|$ and $\langle\phi_x|\phi_y\rangle = \delta_{xy}$. This simple projection measurement is called *von Neumann measurement*. The result x of a von Neumann measurement will occur with probability $\Pr[x] = \langle\psi|E_x|\psi\rangle = |\langle\psi|\phi_x\rangle|^2$. Consider our qubit being represented by the polarization states of a photon. We denote horizontal polarization by $|0\rangle$ and vertical polarization by $|1\rangle$. It is a physical property of the electromagnetic field that these two states are orthogonal,[6] that is, $\langle0|1\rangle = 0$, and thus form a basis in the two-dimensional Hilbert space. A simple measurement that tells us whether the qubit is in the state $|0\rangle$ or $|1\rangle$ is given by the projection set $\{|0\rangle\langle0|, |1\rangle\langle1|\}$. When performing this measurement on the state defined by (1), the qubit will be found in the state $|0\rangle$ with probability $|\alpha|^2$, and in the state $|1\rangle$ with probability $|\beta|^2$. We are free to choose a different basis in the Hilbert space; for instance, the one given by the two states

$$|0\rangle_\times = \frac{1}{\sqrt{2}}(|0\rangle + |1\rangle) \quad \text{and}$$

$$|1\rangle_\times = \frac{1}{\sqrt{2}}(|0\rangle - |1\rangle).$$

This is a rotated basis, and a photon in the state $|0\rangle_\times$ and $|1\rangle_\times$, respectively, has a polarization of $\pm 45°$ against the horizontal. If we measure in this basis by means of the projection measurement $\{|0\rangle_\times\langle0|, |1\rangle_\times\langle1|\}$, we find the qubit in the state $|0\rangle_\times$ with probability $1/2 + \Re(\alpha\bar{\beta})$, and in the state $|1\rangle_\times$ with probability $1/2 - \Re(\alpha\bar{\beta})$. Let us consider the special case where, for instance, $\beta = 0$: When we do the first measurement, we find the qubit in the state $|0\rangle$ with certainty. But when we apply the second measurement, the outcome will be completely random. This is an important property of the *conjugated bases* $\{|0\rangle, |1\rangle\}$ and $\{|0\rangle_\times, |1\rangle_\times\}$ with $|\langle i|j\rangle_\times| = 1/\sqrt{2}$ for all i and j, which will be exploited in many quantum key distribution protocols, as described below.

From POVMs, it is just a small step to *observables*. Each measurable physical quantity is represented by a hermitian operator, called observable. When we write an observable A in its spectral decomposition, $A = \sum_i \lambda_i |i\rangle\langle i|$, where $\langle i|j\rangle = \delta_{ij}$, the corresponding

---

[5]Light particles, called photons, can be seen as electromagnetical waves. A specific property of them is their transversality, which means that the electric and the magnetic fields are orthogonal to each other and to the propagation direction. The inclination of the electric (or magnetic) field to the axis of the propagation is called polarization.

[6]This is by no means a consequence of the geometric relationship between "horizontal" and "vertical." For instance, the spin of a spin-1/2 particle like the electron can point "up" or "down," and the corresponding states $|\uparrow\rangle$ and $|\downarrow\rangle$ are orthogonal. However, the angle between the two spin settings is certainly not 90 degrees.

POVM is given by the orthogonal projectors $\{|i\rangle\langle i|\}$. A measurement of A always yields one of the eigenvalues λ_i as a result, and the measured quantum state *collapses* onto the corresponding state $|i\rangle$.

An important concept in quantum mechanics is the *density matrix* or *density operator* ρ: The density matrix of a so-called *pure* state $|\psi\rangle$ is given by the projector $|\psi\rangle\langle\psi|$. In the case of a qubit, this is a complex-valued (2×2) matrix. The advantage of this representation is the possibility to describe systems with a statistical distribution of states. For instance, consider a system that is known to be in the state $|\psi_x\rangle$ with probability $\Pr[x]$, for $x \in \mathcal{X}$. Let $\mathcal{E} = \{E_y\}_{y \in \mathcal{Y}}$ be some POVM. Then, the probability to get the result y *if* the system was known to be in the state $|\psi_x\rangle$ would be $\langle\psi_x|E_y|\psi_x\rangle$. But since we do not know, we have to average over all possible states, just as we would do if the system were prepared many times in one of the states $\{|\psi_x\rangle\}$ and we had repeated the measurement each time. The probability to measure y in the *ensemble* $\{|\psi_x\rangle, \Pr[x]\}$ is consequently

$$\Pr[y] = \sum_{x \in \mathcal{X}} \Pr[x]\langle\psi_x|E_y|\psi_x\rangle = \mathrm{tr}\left(E_y \sum_{x \in \mathcal{X}} \Pr[x]|\psi_x\rangle\langle\psi_x|\right), \tag{2}$$

where $\mathrm{tr}A$ denotes the trace of the matrix A, that is, the sum of its diagonal elements. We can now introduce the density matrix $\rho = \sum_{x \in \mathcal{X}} \Pr[x]|\psi_x\rangle\langle\psi_x|$, such that (2) takes the simple form: $\Pr[y] = \mathrm{tr}(E_y\rho)$. From now on, we can concentrate on density matrices solely, since any *pure* state $|\psi\rangle$ is just a special case where one probability in the ensemble $\{|\psi_x\rangle, \Pr[x]\}$ is equal to one and all others vanish. In the general case, that is, when at least two different states in the ensemble occur with nonvanishing probability, the system is said to be in a *mixed* state.

Once we consider composite quantum systems, the situation becomes more complicated—and more interesting. Let us consider Alice holding a state ρ_A, acting on a Hilbert space H_A, and Bob holding a state ρ_B acting on H_B. Both states are part of a total state ρ_{AB}, acting on the tensor product $H_A \otimes H_B$, and they are related by the partial trace: $\rho_A = \mathrm{tr}_B\rho_{AB}$ and $\rho_B = \mathrm{tr}_A\rho_{AB}$. This operation discards degrees of freedom in the respective subsystem. Composite states, such as ρ_{AB} can be divided into two classes: *separable* and *entangled* states. We first look at pure states, which means that ρ_{AB} is of the form $\rho_{AB} = |\psi_{AB}\rangle\langle\psi_{AB}|$. Separable pure states are *product states*:

$$|\psi_{AB}\rangle = |\psi_A\rangle \otimes |\psi_B\rangle \equiv |\psi_A\rangle|\psi_B\rangle \equiv |\psi_A\psi_B\rangle.$$

(The last three expressions are equivalent notations.) They are composed of two independent states of the two subsystems A and B. Pure states that cannot be written in this form are called entangled. A famous example of pure entangled states are the Bell states:

$$|\phi^{\pm}\rangle = \frac{1}{\sqrt{2}}(|00\rangle \pm |11\rangle), \tag{3}$$

$$|\psi^{\pm}\rangle = \frac{1}{\sqrt{2}}(|01\rangle \pm |10\rangle). \tag{4}$$

These four states form a basis in the two-qubit Hilbert space. A mixed state is called *separable* if and only if it can be written as a convex sum of projectors onto product

states [Werner 1989]:

$$\rho = \sum_{x \in \mathcal{X}} \Pr[x] |\psi_x^A \phi_x^B\rangle\langle\psi_x^A \phi_x^B| = \sum_{x \in \mathcal{X}} \Pr[x] |\psi_x^A\rangle\langle\psi_x^A| \otimes |\phi_x^B\rangle\langle\phi_x^B|, \qquad (5)$$

with $\Pr[x] \geq 0$ for each $x \in \mathcal{X}$ and $\sum_{x \in \mathcal{X}} \Pr[x] = 1$. These states can be prepared locally in Alice's and Bob's laboratory only by means of *classical* communication, that is, no quantum systems need to be sent. If a state cannot be written in the form (5), it is called *entangled*.

4. QUANTUM KEY DISTRIBUTION

Quantum cryptography exploits the quantum mechanical property that a qubit cannot be copied or amplified without disturbing its original state. This is the statement of the *No-Cloning Theorem* [Wootters and Zurek 1982], which is easily proven: Assume there exists a unitary transformation[7] U that can copy two states $|\psi_1\rangle$ and $|\psi_2\rangle$:

$$U|\psi_1\rangle|0\rangle = |\psi_1\rangle|\psi_1\rangle, \qquad (6)$$
$$U|\psi_2\rangle|0\rangle = |\psi_2\rangle|\psi_2\rangle, \qquad (7)$$

where $|0\rangle$ is an arbitrary input state. If we equate the scalar products of the left-hand and right-hand sides, it follows by the unitarity of U that $\langle\psi_1|\psi_2\rangle = \langle\psi_1|\psi_2\rangle^2$, which implies that $\langle\psi_1|\psi_2\rangle$ equals 0 or 1. This means that we can copy only orthogonal or identical states. In contrast, arbitrary unknown states cannot be perfectly cloned. (Note that orthogonal or identical states are not viewed as "unknown" states, since we do know they are orthogonal, for example.)

The essence of this theorem is the main ingredient of quantum key distribution, where Alice and Bob use a quantum channel to exchange a sequence of qubits, which will then be used to create a key for the one-time pad in order to communicate over an insecure channel. Any disturbance of the qubits, for example caused by Eve trying to measure the qubits' state, can be detected with high probability.

In this section, we describe the BB84 protocol proposed by Bennett and Brassard [1984].[8] This is the first protocol designed to employ quantum mechanics for two parties to agree on a joint secret key.

4.1. The BB84 Protocol

In this protocol, Alice and Bob use a quantum channel to send qubits. They are also connected by a classical channel, which is insecure against an eavesdropper but unjammable. Alice and Bob use four possible quantum states in two conjugate bases (say, the rectilinear basis $+$ and the diagonal basis $\times$). We use $|0\rangle_+$ and $|0\rangle_\times = (|0\rangle_+ + |1\rangle_+)/\sqrt{2}$ for the classical signal "0," and we use $|1\rangle_+$ and $|1\rangle_\times = (|0\rangle_+ - |1\rangle_+)/\sqrt{2}$ for the classical signal "1." Note that the two bases are connected by the so-called Hadamard transformation

$$H = \frac{1}{\sqrt{2}} \begin{pmatrix} 1 & 1 \\ 1 & -1 \end{pmatrix} \qquad (8)$$

in the following way: We have $H|0\rangle_+ = |0\rangle_\times$ and $H|1\rangle_+ = |1\rangle_\times$, and vice-versa, since $H^2 = \mathbb{1}$.

---

[7]The time-evolution of an isolated quantum system is described by a unitary transformation $U\colon |\psi\rangle \to U|\psi\rangle$.
[8]Some of the ideas used in the BB84 protocol were already introduced by Wiesner [1983].

Table I. The BB84 Key Distribution Protocol. Here, "Y" and "N" stand for "yes" and "no," respectively, and "R" means that Bob obtains a random result

Alice's string	1	1	0	1	0	0	1	0	1	1	1	1	0	0
Alice's basis	+	+	+	×	×	+	×	×	×	×	+	+	+	+
Bob's basis	+	×	+	+	×	+	×	+	×	×	+	+	+	+
Bob's string	1	R	0	R	0	0	1	R	1	1	1	1	0	0
Same basis?	Y	N	Y	N	Y	Y	Y	N	Y	Y	Y	Y	Y	Y
Bits to keep	1		0		0	0	1		1	1	1	1	0	0
Test	Y		N		N	Y	N		N	N	N	Y	Y	N
Key			0		0		1		1	1	1			0

The protocol works as follows (see also Table I for illustration):

(1) Alice randomly prepares $2n$ qubits, each in one of the four states $|0\rangle_+$, $|0\rangle_\times$, $|1\rangle_+$, or $|1\rangle_\times$, and sends them to Bob.

(2) For each qubit that Bob receives, he chooses at random one of the two bases ($+$ or $\times$) and measures the qubit with respect to that basis. In the case of a perfectly noiseless channel, if Bob chooses the same basis as Alice, his measurement result is the same as the classical bit that Alice prepared. If the bases differ, Bob's result is completely random.

(3) Alice tells Bob via the classical channel which basis she used for each qubit. They keep the bits where Bob has used the same basis for his measurement as Alice. This happens in about half the cases, so they will have approximately n bits left. These are forming the so-called *sifted key*.

(4) Alice and Bob choose a subset of the sifted key to estimate the error-rate. They do so by announcing publicly the bit values of the subset. If they differ in too many cases, they abort the protocol, since its security cannot be guaranteed.

(5) Finally, Alice and Bob obtain a joint secret key from the remaining bits by performing error correction and privacy amplification.

Which possibilities does Eve have to attack this protocol? And, consequently, what is the threshold of the error-rate, at which Alice and Bob should abort the protocol? To answer these questions, we look at a simple eavesdropping strategy, which is called "intercept-and-resend." (This attack is not the optimal one from Eve's perspective, that is, there are strategies that provide the adversary with more information about the key.) We will not rigorously prove the security of the protocol against the "intercept-and-resend" attack here. Rather, we consider this attack merely to provide some intuition about how the BB84 protocol counteracts eavesdropping.

Eve's goal is to learn at least some part of the key. Thus, an obvious strategy for her is to intercept the qubits being transmitted from Alice to Bob. She cannot simply copy the qubits, since this would contradict the No-Cloning Theorem. In order to extract some information, she is forced to measure (and thus destroy) them. But since she does not know the basis in which they were prepared (Alice announces this information only after Bob received all signals), she can only guess or just flip a coin for the selection of the measurement basis. In about half the cases, she will happen to choose the same basis as Alice and get completely correlated bit values. In the other half, her results will be random and uncorrelated. Bob certainly expects to receive something from Alice, so Eve needs to send some qubits to him. However, she still has no idea which basis Alice used, so she prepares each qubit in the same basis as she measured it (or she chooses a basis at random). These newly created qubits again match Alice's bases in only half of the cases. After Bob receives Eve's qubits, he measures them, and Alice and Bob apply the sifting. Because of Eve's disturbance, about half of Bob's key was measured in a different basis than it was prepared by Alice. Since Bob's result is random in those

cases, his sifted key will contain about 25% errors. In the error-estimation stage, if Alice and Bob obtain such a high error rate, it would be wise for them to abort the protocol.

If the error rate is below an agreed threshold value, Alice and Bob can eliminate errors with (classical) error correction. A simple method for error correction works as follows: Alice chooses two bits at random and tells Bob the XOR-value of the two bits. Bob tells Alice if he has the same value. In this case, they keep the first bit and discard the second bit. If their values differ, they discard both bits. The remaining bits form the key.

The last stage of the protocol is privacy amplification [Maurer 1993; Bennett et al. 1995]—a procedure in which Alice and Bob eliminate (or, at least, drastically reduce) Eve's knowledge about the key. They do so by choosing random pairs of bits of the sifted key and replacing them by their corresponding XOR-values. Thus, they halve the length of the key, in order to "amplify" their privacy. Note that Eve has less knowledge about the XOR-value, even if she knew the values of the single bits with high probability (but not with certainty).

Note that these simple methods for error correction and privacy amplification do not always work. For the general case, there exist more sophisticated strategies. For more details on error-correcting codes and their usage in the physics of quantum information, we refer to Huffman and Pless [2003] and Bouwmeester et al. [2000].

4.2. Security of Quantum Key Distribution

Unlike many of the classical cryptosystems in use today, whose security often draws on unproven assumptions about the computational complexity of mathematical problems, the security of quantum cryptography is based on—and employs—the laws of physics. The term "unconditional security" is used to emphasize the fact that it does not rely on the presumed, yet unproven hardness of some mathematical problem. In this section, we present the proof of the unconditional security of the BB84 protocol, as devised by Shor and Preskill [2000].

We divide the proof into three parts:

—In the first part, we present the so-called *entanglement-based* version of the BB84 protocol. In contrast, the scheme presented in the previous section is called a *prepare-and-measure scheme*, for obvious reasons. In the entanglement-based version, Alice and Bob's aim is to share a special entangled state that allows them to obtain perfectly correlated bits upon measuring their half of the state. We will see how they can construct such a state, how they can check whether they were successful, and how they can detect Eve's attempted attack.

—In the second part, we will show that the equivalent entanglement-based version is secure. In contrast to earlier work by Shor and Preskill [2000], which is based on a proof by Lo and Chau [1999], we use the *universally composable* definition of unconditional security [Ben-Or et al. 2005]. This definition ensures that the key can be used in any composition of cryptographic primitives.

—In the third part, we show that the two schemes are equivalent indeed.

4.2.1. The Entanglement-Based Version of BB84. The entanglement-based version of the BB84 protocol that we now present is similar to the protocol introduced by Ekert [1991] and follows ideas of Bennett et al. [1992]. In this version of the protocol, Alice and Bob aim at creating a special entangled state, namely the Bell state

$$|\phi^+\rangle = \frac{1}{\sqrt{2}}(|00\rangle + |11\rangle), \tag{9}$$

where Alice holds the first particle and Bob holds the second one. An important property of this state is that it has the same form in the rectilinear basis $+$ and in the diagonal basis $\times$, as $|\phi^+\rangle = \frac{1}{\sqrt{2}}(|0\rangle_+|0\rangle_+ + |1\rangle_+|1\rangle_+) = \frac{1}{\sqrt{2}}(|0\rangle_\times|0\rangle_\times + |1\rangle_\times|1\rangle_\times)$. This means that Alice's and Bob's measurement results are completely correlated whenever they measure the state $|\phi^+\rangle$ in one of those bases. (Moreover, their results are random.) Since the state is pure, it cannot be entangled with anything else, in particular not with anything under Eve's control. Thus, whenever Alice and Bob are sure they share a $|\phi^+\rangle$ state, they know that (a) measuring in the same basis generates a shared random bit, and (b) Eve has no knowledge about this bit. To generate the whole key, Alice and Bob prepare a large number of these Bell states,

$$|\phi^+\rangle^{\otimes n} = |\phi^+\rangle \otimes \cdots \otimes |\phi^+\rangle,$$

and measure each qubit separately. We will now show how they can achieve this.

We need to take a brief detour to quantum error correction first. In contrast to a classical bit, a qubit can undergo three different errors: bit flips, phase errors, and combinations of these two:

—When a bit flip occurs, the state $|0\rangle$ becomes $|1\rangle$, and vice-versa. This error is described by the Pauli matrix

$$\sigma_x = \begin{pmatrix} 0 & 1 \\ 1 & 0 \end{pmatrix}.$$

—Phase errors transform the state $|1\rangle$ into $-|1\rangle$, but leave $|0\rangle$ unchanged. Such an error is described by the Pauli matrix

$$\sigma_z = \begin{pmatrix} 1 & 0 \\ 0 & -1 \end{pmatrix}.$$

—Both these errors can also occur combined. For example, changing $|0\rangle$ to $-|1\rangle$ and $|1\rangle$ to $|0\rangle$ can be described by $\sigma_z\sigma_x = i\sigma_y$, where

$$\sigma_y = \begin{pmatrix} 0 & -i \\ i & 0 \end{pmatrix}.$$

Let us now recall some elements of classical error correction. A (classical) linear $[n, k]$ code C that encodes k bits of information by an n bit string is a set of 2^k codewords. Each codeword is an n-dimensional binary vector. The whole code can be described by an $(n \times k)$-dimensional generator matrix G that maps each message x to the encoded message Gx. Thus, the set of all possible codewords is the vector space that is spanned by the columns of G. We require those vectors to be linearly independent. Error correction for linear codes can be easily described by means of the parity check matrix H. This is an $((n - k) \times n)$ matrix with the property that $Hx = 0$ for all codewords x.

Suppose now that a message x is encoded as $y = Gx$. Due to an error e, one obtains $y' = y + e$. Since we have $Hy = 0$ for all codewords, it follows that $Hy' = He$, which is called the *(error) syndrome*. Thus, if the syndrome is 0, no error has occurred. Otherwise, H is constructed such that the syndrome contains information about the error that should make it possible to correct it. Finally, we introduce the concept of *duality*: Let C be a linear $[n, k]$ code with generator matrix G and parity check matrix H. Then, we can define the dual code $C^\perp$ of C, which is the set of all codewords that are orthogonal to each codeword in C. The dual code $C^\perp$ is an $[n - k, n]$ code which is generated by H^T and has a parity check matrix G^T. Dual codes play an important role in the construction of CSS codes, as we explain below.

Definition 4.1. Let C_1 and C_2 be classical linear $[n, k_1]$ and $[n, k_2]$ codes, respectively, such that $C_2 \subset C_1$. For each codeword $x \in C_1$, define the quantum state

$$|x + C_2\rangle = \frac{1}{\sqrt{|C_2|}} \sum_{y \in C_2} |x + y\rangle. \tag{10}$$

The space spanned by $\{|x + C_2\rangle\}_{x \in C_1}$ defines an $[n, k_1 - k_2]$ quantum code, which is called the *Calderbank–Shor–Steane code*, CSS(C_1, C_2) for short.

Let x and x' in C_1 be codewords such that $x - x'$ is in C_2. Then one can show that

$$|x + C_2\rangle = |x' + C_2\rangle,$$

that is, the state $|x + C_2\rangle$ depends only on C_1/C_2, that is, on the coset to which x belongs.[9] It follows that if x and x' belong to different cosets, the states $|x + C_2\rangle$ and $|x' + C_2\rangle$ are orthogonal. As the number of cosets of C_2 in C_1 is $|C_1|/|C_2|$, the dimension of the space CSS(C_1, C_2) is $|C_1|/|C_2| = 2^{k_1 - k_2}$, thus $m = k_1 - k_2$ qubits can be encoded.

Error correction with CSS codes works as follows. Suppose that C_1 and $C_2^\perp$ both can correct ℓ errors. Moreover, let H_1 be the parity check matrix for C_1, and let H_2 be that for $C_2^\perp$. Define

$$\sigma_\alpha^s = \sigma_\alpha^{s_1} \otimes \sigma_\alpha^{s_2} \otimes \cdots \otimes \sigma_\alpha^{s_n}, \tag{11}$$

where $\alpha \in \{x, y, z\}$, $\sigma_\alpha^0 = \mathbb{1}$, and $s = (s_1, s_2, \ldots, s_n)$ is an n bit vector. It can be shown that the syndrome for bit flip errors can be computed by measuring σ_z^r for each row vector r of H_1. Similarly, the syndrome for phase errors can be computed by measuring σ_x^t for each row vector t of H_2. In this way, ℓ bit flips and ℓ phase errors can be corrected.

We have now collected all the ingredients to describe the entanglement-based version of the BB84 protocol:

(1) Alice creates $2n$ qubit pairs in the state $|\phi^+\rangle^{\otimes 2n}$.
(2) She randomly selects n of those qubits which will later serve as check qubits.
(3) Alice selects a random $2n$ bit string b and applies the Hadamard transformation (8) to her half of each qubit pair whenever the corresponding bit of b is "1."
(4) She sends the other half of all qubit pairs to Bob.
(5) Alice announces b and which qubits are to serve as check qubits.
(6) Bob performs a Hadamard transformation on those of his qubits where b is "1."
(7) Alice and Bob measure the check qubits in the $\{|0\rangle, |1\rangle\}$ basis to estimate the error rate. If more than ℓ results differ, they abort the protocol.
(8) For the remaining qubits, Alice and Bob measure the syndromes for the codes C_1 and C_2, correct the errors, and obtain $|\phi^+\rangle^{\otimes m}$.
(9) They measure this state in the $\{|0\rangle, |1\rangle\}$ basis to obtain a shared secret key.

The point of performing the Hadamard transformation on half of the qubits is that this operation effectively changes the basis, in which the qubits are prepared, from $\{|0\rangle_+, |1\rangle_+\}$ to $\{|0\rangle_\times, |1\rangle_\times\}$. This is necessary because if Eve knew the basis, she could launch the intercept-resend attack presented in the previous section and break the protocol.

---

[9]Let G and H be two groups with $G \subset H$. For each $h \in H$, we define the *(left) coset of G in H with respect to h* as $hG = \{h + g \mid g \in G\}$. The group H/G is the set of all cosets of G in H (i.e., the equivalence classes).

4.2.2. Security of the Entanglement-Based Version. Up to this point, we often used the term "security" without providing a rigorous definition. In this section, we will make up for this. Additionally, we need to provide a mathematical framework to cover *all* possible eavesdropping strategies, in particular those where the adversary stores a quantum system that contains information about the classical bit strings obtained by Alice and Bob upon measuring their quantum states. Such a situation, where a quantum system is correlated with classical data, can be described by so-called *classical-quantum states* (cq-states, for short): Let X be a random variable with range $\mathcal{X}$ and let $\{|x\rangle\}_{x\in\mathcal{X}}$ be some basis of a Hilbert space. Moreover, denote by ρ_E^x the state of the quantum system E conditioned on the value x of the random variable X. Then, the overall system can be described by the cq-state

$$\rho_{XE} = \sum_{x\in\mathcal{X}} \Pr[x]|x\rangle\langle x| \otimes \rho_E^x.$$

Applying this formalism to our key distribution scenario, let $\mathcal{S}$ denote the set of all possible keys that can be extracted by the protocol. The individual keys held by Alice and Bob can be described by random variables S_A and S_B, respectively, taking values s_A and s_B in $\mathcal{S}$. The adversary holds a quantum system $\rho_E^{s_A s_B}$, which is correlated with those variables, and thus the total system can be described by the *classical-classical-quantum state* (ccq-state, for short)

$$\rho_{S_A S_B E} = \sum_{s_A, s_B \in \mathcal{S}} \Pr[s_A, s_B]|s_A\rangle\langle s_A| \otimes |s_B\rangle\langle s_B| \otimes \rho_E^{s_A s_B}. \tag{12}$$

In the ideal case, Alice's and Bob's keys are identical and uniformly distributed, that is, each possible key occurs with equal probability. Moreover, the state of Eve's quantum system should be completely independent of the key. Thus, the ideal ccq-state is given by

$$\rho_{UU} \otimes \rho_E = \frac{1}{|\mathcal{S}|} \sum_{s\in\mathcal{S}} |s\rangle\langle s| \otimes |s\rangle\langle s| \otimes \rho_E. \tag{13}$$

We now are ready to define the notion of unconditionally secure key.

Definition 4.2.   Let $\rho_{S_A S_B E}$, as defined in (12), be the ccq-state describing a classical key pair (S_A, S_B) together with an adversary holding a quantum system E. Then (S_A, S_B) is said to be ϵ-*secure with respect to* E if and only if

$$\|\rho_{S_A S_B E} - \rho_{UU} \otimes \rho_E\| \leq \epsilon,$$

where $\rho_{UU} \otimes \rho_E$ is the ideal state, defined in (13).

Here, $\|\rho - \sigma\| = \mathrm{tr}|\rho - \sigma|/2$ (with $|M| = \sqrt{M^\dagger M}$ where $M^\dagger = \overline{M}^T$ is the conjugate transpose of matrix M) denotes the trace distance, which is a proper distance measure in the space of hermitian operators. The above definition of security (Definition 4.2.) has the intuitive interpretation that a key pair (S_A, S_B) is ϵ-secure if it is "ϵ-close" to the ideal state described by (13) in the sense that (S_A, S_B) is an ideal pair of keys with probability at least $1 - \epsilon$. Moreover, it guarantees that the key pair remains secure when used in cryptographic applications.

We can now prove the unconditional security of the entanglement-based version of the BB84 protocol. Recall that the aim of this protocol is to distribute the state $|\phi^+\rangle^{\otimes n}$. In the real world, Alice and Bob are of course not able to *exactly* achieve this; rather,

at the end of the protocol (after Step (8)), they will hold a state ρ_{AB}, which hopefully is very similar to $|\phi^+\rangle^{\otimes n}$. The "distance" to a pure state is measured by means of the so-called *fidelity*, which is defined as $F(\rho, |\psi\rangle) = \langle\psi|\rho|\psi\rangle$. If $F = 1$, the two states are identical. Since we do not make any restrictions about the eavesdropper's strategy, we consider the worst case in which Eve holds a purifying system of ρ_{AB}. This is the state $\rho_E = \text{tr}_{AB}|\Psi_{ABE}\rangle\langle\Psi_{ABE}|$, where $|\Psi_{ABE}\rangle$ is a pure state (in a higher-dimensional Hilbert space) such that $\rho_{AB} = \text{tr}_E|\Psi_{ABE}\rangle\langle\Psi_{ABE}|$. This scenario corresponds to the case where the adversary has full control over the quantum channel.

The following lemma relates the fidelity of ρ_{AB} to $|\phi^+\rangle^{\otimes n}$ with the security of the key that is obtained when measuring ρ_{AB}. The proof of this lemma can be found in König et al. [2006].

LEMMA 4.3. *Let $\epsilon \geq 0$ and ρ_{AB} be a bipartite quantum state such that*

$$F(\rho_{AB}, |\phi^+\rangle^{\otimes n}) \geq \sqrt{1 - \epsilon^2}.$$

Then, the two n bit strings obtained from measuring ρ_{AB} locally in the $\{|0\rangle, |1\rangle\}$-basis are ϵ-secure keys, with respect to an adversary holding the purifying system of ρ_{AB}.

It remains to show that by the random sampling that Alice and Bob apply, they can reliably estimate the fidelity of the remaining qubits. The main ingredient to prove this is again a lemma, which we also state here without proof. (The proof is left to the reader, see Nielsen and Chuang [2000].)

LEMMA 4.4. *Let a random 2n bit string that might contain some errors, and a random subset of n check bits of that string be given. Then, for any two constants $\delta > 0$ and $\epsilon > 0$, the probability of finding less than δn errors on the check bits, and more than $(\delta + \epsilon)n$ errors on the remaining bits is less than $e^{-\mathcal{O}(\epsilon^2 n)}$, for sufficiently large n.*

Although this lemma is based on classical probability theory, we can give an argument for its validity in the quantum world: The observables that Alice and Bob measure on the check bits are both diagonal in the Bell basis (Eqs. (3) and (4)), which means that the statistics of the results can be described purely classically. These measurements on $H_A \otimes H_B$ are given by the POVMs

$$\{P_{\text{bf}} = |\psi^+\rangle\langle\psi^+| + |\psi^-\rangle\langle\psi^-|, \mathbb{1} - P_{\text{bf}}\},$$

which are used to check for bit flips, and

$$\{P_{\text{pe}} = |\phi^-\rangle\langle\phi^-| + |\psi^-\rangle\langle\psi^-|, \mathbb{1} - P_{\text{pe}}\},$$

which are used to check for phase errors. Alice and Bob choose one of those measurements at random for each check qubit. In this way, they can calculate a lower bound for the fidelity of the remaining qubits.

To summarize, we have shown that by random sampling the fidelity of the state shared by Alice and Bob can be lower-bounded, with an exponentially small probability of error. Moreover, this bound directly defines how secure a key generated by measuring this state will be.

4.2.3. Equivalence of the Two Schemes. We prove the equivalence of the entanglement-based and prepare-and-measure versions of the BB84 protocol by successive simplifications. Each step is very simple, so it is easy to verify that the security of the protocol is not compromised.

A major simplification is that all measurements done by Alice after transmitting the particles can already be done at the very beginning: If Alice measures her part of the state $|\phi^+\rangle$, she obtains a random bit as a result, but on the other hand, Bob's part of the state collapses onto the correlated state $|0\rangle$ or $|1\rangle$. Thus, instead of sending entangled qubits for the check, Alice can as well prepare single qubits randomly in one of the states $|0\rangle$ and $|1\rangle$, and send those states to Bob. Of course, it is crucial for the security of the protocol that Eve does not know *a priori* which qubits will serve as check qubits and which as "key qubits"; otherwise, she could treat them differently and thus fudge the error estimation.

Another measurement Alice can do at the beginning is the measurement of her syndrome and her key qubits. This is not very obvious, so let us give some more detail: Given a CSS code $\mathrm{CSS}(C_1, C_2)$, we can define a family of equivalent codes $\mathrm{CSS}_{v,w}(C_1, C_2)$, in the sense that they have the same error correcting properties. The codewords of the code $\mathrm{CSS}_{v,w}(C_1, C_2)$ are given by

$$|x_k, v, w\rangle = \frac{1}{\sqrt{|C_2|}} \sum_{y \in C_2} (-1)^{v \cdot y} |x_k + y + w\rangle, \qquad (14)$$

where x_k is one representative of one of the m cosets of C_2 in C_1, and v and w are arbitrary n bit strings. Since the $\{|x_k, v, w\rangle\}$ form a basis, we can rewrite

$$|\phi^+\rangle^{\otimes n} = \frac{1}{\sqrt{2^n}} \sum_{i=0}^{2^n-1} |i\rangle|i\rangle = \frac{1}{\sqrt{2^n}} \sum_{x_k, v, w} |x_k, v, w\rangle|x_k, v, w\rangle, \qquad (15)$$

where i is in binary notation. If now Alice measures the error syndromes, namely σ_z^r for each row vector r of H_1 and σ_x^t for each row vector t of H_2, she obtains a random result for v and w. Finally, if she does a last measurement in the $\{|0\rangle, |1\rangle\}$ basis, she obtains a random codeword x_k. From (15), we see that Bob's state then collapses onto $|x_k, v, w\rangle$, which is a random qubit encoded in a random code.

As an intermediate result, we rephrase the entanglement-based protocol including all simplifications introduced so far:

(1) Alice creates n random check qubits, each in the state $|0\rangle$ or $|1\rangle$, a random n bit string k, which will serve as the key, and two random n bit strings v and w. She prepares the state $|k\rangle$ and encodes it using $\mathrm{CSS}_{v,w}(C_1, C_2)$.
(2) She randomly selects n positions for the check qubits and puts the encoded qubits in the remaining positions.
(3) Alice selects a random $2n$ bit string b and applies the Hadamard transformation to her half of each qubit pair where b is "1."
(4) She sends the other half of all qubit pairs to Bob.
(5) Alice announces b, v, and w, and which qubits are to serve as check qubits.
(6) Bob performs a Hadamard transformation on those of his qubits where b is "1."
(7) Bob measures the check qubits in the $\{|0\rangle, |1\rangle\}$ basis. If he finds more than ℓ results that disagree with Alice's prepared states, they abort the protocol.
(8) Bob decodes the key qubits from $\mathrm{CSS}_{v,w}(C_1, C_2)$ and obtains the state $|k\rangle$.
(9) He measures $|k\rangle$ in the $\{|0\rangle, |1\rangle\}$ basis and obtains the key k as the result.

We will now simplify this protocol even further: Note that in the original version, Alice and Bob do not care whether they shared the state $|\phi^+\rangle$ or $|\phi^-\rangle = (|00\rangle - |11\rangle)/\sqrt{2}$,

because measuring both states provides them with correlated, random bits; the relative phase is irrelevant. Thus, it is unnecessary to send the phase correction information v to Bob. This is why CSS codes are used: They decouple the bit flip error correction from the phase error correction. If now Bob were to measure his key qubits before the decoding, he would obtain $x_k + y + w + e$, where e denotes the bit errors that occurred during the transmission (or that were introduced by Eve). He can now classically decode this bit string by subtracting w, which was announced by Alice, and correct it to the codeword $x_k + y$, if e did not introduce too many errors. Bob finds the key by computing the coset to which $x_k + y$ belongs. But since Bob does not need v, why should Alice send it? If she never reveals that value, she effectively prepares a state that is a classical mixture of all possible values that v can take, weighted with the corresponding probabilities:

$$\rho_{x_k, w} = \frac{1}{2^n} \sum_v |x_k, v, w\rangle\langle x_k, v, w| = \frac{1}{|C_2|} \sum_{z \in C_2} |x_k + z + w\rangle\langle x_k + z + w|. \qquad (16)$$

We see that this state can also be prepared by classically choosing a random codeword $z \in C_2$ and constructing $|x_k + z + w\rangle$. Thus, the preparation in Step (1) can be done equivalently in the following way: Alice creates n random check qubits, each in the state $|0\rangle$ or $|1\rangle$, a random n bit string w, a random string $x_k \in C_1/C_2$, and a random codeword $z \in C_2$. The n key qubits are prepared in the state $|x_k + w + z\rangle$, and the check qubits are placed at random positions.

Note that we can also remove the need for $z \in C_2$, if Alice instead of choosing $x_k \in C_1/C_2$ chooses $x_k \in C_1$. With this modification, Alice sends the state $|x_k + w\rangle$ as key qubits, which Bob then measures and corrects to $x_k + w$. Since $x_k + w$ is a completely random n bit string, Alice can as well just prepare $|y\rangle$, where y is a random n bit string. She sends it to Bob who measures it to obtain $y + e$, then Alice sends error correction information $y - x_k$, which Bob subtracts from $y + e$ to finally obtain $x_k + e$. He corrects it to x_k and calculates the key k as the coset to which x_k belongs. What we have achieved is that now the check and the key qubits are just prepared randomly in one of the states $|0\rangle$ or $|1\rangle$. The whole protocol so far looks as follows:

(1) Alice creates $2n$ random qubits, each in the state $|0\rangle$ or $|1\rangle$, and a random codeword $x_k \in C_1$.
(2) She randomly selects n positions to be check qubits and the remaining n positions to define the key qubits $|y\rangle$.
(3) Alice selects a random $2n$ bit string b and applies the Hadamard transformation to her half of each qubit pair where b is "1."
(4) She sends the other half of all qubit pairs to Bob.
(5) Alice announces b and $y - x_k$, and which qubits are to serve as check qubits.
(6) Bob performs a Hadamard transformation on those of his qubits where b is "1."
(7) Bob measures the check qubits in the $\{|0\rangle, |1\rangle\}$ basis. If he finds more than ℓ results that disagree with Alice's prepared state, they abort the protocol.
(8) Bob measures the key qubits and gets $y + e$, subtracts $y - x_k$, and corrects $x_k + e$ to x_k.
(9) He calculates the coset to which x_k belongs to get the key k.

Finally, we can remove the Hadamard transformation, and let Alice choose randomly one of the four states in $\{|0\rangle_+, |1\rangle_+, |0\rangle_\times, |1\rangle_\times\}$. Then Bob, instead of waiting for b to be announced, simply chooses one basis at random and measures the arriving qubits. As he will choose the wrong basis in roughly half the cases, Alice should double the number

of input qubits to $4n$. After his measurement, Alice announces which basis she used and both discard all instances where they used a different basis. With this last modification, we finally arrived at the prepare-and-measure version of the BB84 protocol, only up to some small twists.

5. QUANTUM BIT COMMITMENT

When talking about quantum cryptography, everyone is thinking about key distribution. There are, however, other cryptographic applications as well, such as bit commitment. A bit commitment protocol based on quantum mechanics was introduced by Brassard et al. [1993]. The unconditional security of the protocol (which means that the security of the protocol is independent of the computational resources, such as computing time, amount of memory used, and computer technology of the cheater) has been accepted without proof [Yao 1995]. Two years after it had been proposed, the protocol turned out to be insecure [Mayers 1995].

A commitment protocol is a procedure in which one party, say Alice, deposits a message such that no one (and in particular not Alice) can read it nor change it. At some point in the future, Alice can announce her message, and with high certainty it can be proven that the revealed message is the same as the one Alice had deposited originally. To illustrate this situation, suppose Bob wants to auction off a diamond ring, subject to the condition that each person wishing to participate in the auction can bid only one single amount of money. After each person has chosen a specific amount, the highest bidder gets the ring. So everyone writes their own bid on a piece of paper, puts it into a personal safe, which is then locked and given to Bob. Until all bids have been submitted to Bob, each bidder keeps the key matching the lock of his or her safe. In this way Bob cannot see any of the bids, which in turn cannot be changed once they have been submitted, since only Bob has access to the committed safes. All keys are handed over to Bob after he has received all safes from the people participating in the auction. The different offers are compared in public, so that everybody can be sure that only the highest bidder walks away with the diamond and an empty wallet.

We can describe this commitment protocol mathematically as follows: The protocol has two stages, the commit phase and the unveil phase. Alice commits herself to the data m by computing $c = f(m)$, and she sends c to Bob. Alice unveils the commitment by showing Bob the preimage m of c. In classical cryptography, and in particular in public-key cryptography, one-way functions are used for commitment. In quantum cryptography, we want to make use of the laws of quantum mechanics to create a fair protocol for both sides. Bit commitment is a special case of a commitment protocol, where the data m consists of only one single bit.

It is widely believed that it is impossible to create a perfectly secure classical bit commitment protocol. Regarding the extension to the quantum world, it was shown that unconditionally secure quantum bit commitment is also impossible [Mayers 1997; Lo and Chau 1997]. However, when relaxing the security constraints, quantum bit commitment becomes possible in slightly modified frameworks. One example is Kent's quantum bit commitment protocol, which is based on special relativity theory [Kent 1999]. Another example is due to Damgård et al. [2005] who proposed a quantum bit commitment protocol that is secure in the bounded storage model.

5.1. The BB84 Quantum Bit Commitment Protocol

The BB84 protocol was introduced in Section 4.1. A quantum bit commitment protocol can be created from the BB84 quantum key distribution protocol with a few minor

changes [Bennett and Brassard 1984]. Just as in the classical bit commitment protocol, the quantum protocol starts with the commit phase and ends with the unveil phase.

The commit procedure:

(1) Alice chooses a bit $b \in \{0, 1\}$.
(2) Alice creates a random binary string $w = w_1 \cdots w_n$ with n bits.
(3) If Alice wants to commit to 0, she does a quantum encoding of each bit w_i in the two basis states of the rectilinear basis $+$. If she wants to commit to 1, she encodes the bits in the two basis states of the diagonal basis $\times$. Let θ_i denote the basis chosen for w_i.
(4) Alice sends the sequence of n encoded quantum states to Bob.
(5) Bob chooses a random measurement basis (rectilinear or diagonal) for each of the received quantum states, i.e., he chooses a string of random bases $\hat{\theta} = \hat{\theta}_1 \cdots \hat{\theta}_n \in \{+, \times\}^n$. He measures the ith state in the basis $\hat{\theta}_i$, and denotes the outcome by $\hat{w}_i$.

If we take a look at the two density matrices for the n states corresponding to $b = 0$ and $b = 1$, respectively, it is easy to see that they are the same, and equal to the identity matrix. Thus, Bob has no chance to get any information about the bit b.

The unveil procedure:

(1) Alice publishes b (i.e., the basis that she used for encoding) and the string w.
(2) For about half of the n states, Bob used the same basis for his measurement as Alice used for encoding. In these cases Bob can verify that Alice's revealed bits are matching his measurement results.

How could a dishonest party cheat in this protocol? For example, Alice could choose the bit $b = 1$ for the commit phase, so she encodes the states with the diagonal basis $\times$. Later during the unveil phase, she changes her mind and tells Bob that she committed to the bit $b = 0$, so Bob assumes that Alice has used the rectilinear basis $+$. In approximately $n/2$ cases, Bob measures the states with the rectilinear basis $+$, and in these cases Alice has to guess the bits Bob measured. Since Alice's success to make a right guess for one bit is $1/2$, her overall cheating will not be detected with a probability of $(1/2)^{n/2}$. Once n is chosen large enough, Alice has practically no chance to manipulate the protocol by this probabilistic method.

But what if Alice uses specially entangled states as in the entanglement-based version of the BB84 protocol (see Section 4.2.1, Eq. (9))? Alice could create n pairs of entangled states and send one part of each pair to Bob. She doesn't have to commit to a bit in the beginning, because she can perform a measurement right before the unveil phase. If, for example, she chooses bit $b = 0$, she measures the states that she has kept in the rectilinear basis $+$. Bob's measurement results will be perfectly correlated, due to the shape of the entangled state in Eq. (9). If Alice wants to choose bit $b = 1$ instead, she measures the states that she has kept in the diagonal basis $\times$. As the state from Eq. (9) is form-invariant under a basis rotation by $45°$, Alice's announced encoded states will again match Bob's measurement results. Thus, Bob has no chance to notice the attack.

5.2. Impossibility of Unconditionally Secure Quantum Bit Commitment

As mentioned above, unconditionally secure quantum bit commitment is impossible. In this section we will review the main arguments to prove this statement. According to Lo and Chau [1997], the ideas of all quantum bit commitment protocols proposed up to date can be roughly described by the following five steps:

(1) Alice chooses a bit $b \in \{0, 1\}$ and prepares the state

$$|0\rangle = \sum_i \alpha_i |e_i^A\rangle \otimes |f_i^B\rangle$$

for $b = 0$, and the state

$$|1\rangle = \sum_j \beta_j \left|e_j'^A\right\rangle \otimes \left|f_j'^B\right\rangle$$

for $b = 1$, where $|e_i^A\rangle$ and $|e_j'^A\rangle$ are orthonormal bases of Alice's Hilbert space, that is, $\langle e_i^A | e_k^A \rangle = \delta_{ik}$ and $\langle e_j'^A | e_l'^A \rangle = \delta_{jl}$. The states $|f_i^B\rangle$ and $|f_j'^B\rangle$ live in Bob's Hilbert space, and are not necessarily orthogonal to each other.

(2) Now, Alice has to make a measurement on the first part of the above state, and will thus determine i or j, depending on her initial choice for b.

(3) Alice sends the second part of the above state to Bob. This is the last step in the commit phase.

(4) At the beginning of the unveil phase, Alice publicly announces i or j together with b.

(5) Bob makes a measurement on his part of the state, in order to make sure that in Step (3), Alice committed to the same bit she has announced in Step (4).

To show that a cheating Alice cannot be detected, we distinguish two cases. We give only a sketch of the proof, for more details we refer to Mayers [1997] and Lo and Chau [1997].

We first consider the case where Bob cannot get any information about the bit b out of the state that Alice sent him. This means that his two possible reduced density matrices, corresponding to the two states $|0\rangle$ and $|1\rangle$, are the same, that is, $\text{tr}_A|0\rangle\langle 0| = \text{tr}_A|1\rangle\langle 1|$. Now, we can write the Schmidt decomposition (i.e., a bi-orthogonal decomposition that can always be found, see, for example, Nielsen and Chuang [2000]) as

$$|0\rangle = \sum_k \sqrt{\lambda_k} |\hat{e}_k^A\rangle \otimes |\hat{f}_k^B\rangle$$

and

$$|1\rangle = \sum_k \sqrt{\lambda_k} \left|\hat{e}_k'^A\right\rangle \otimes |\hat{f}_k^B\rangle,$$

where $|\hat{e}_k^A\rangle$ and $|\hat{e}_k'^A\rangle$ are orthonormal bases of Alice's Hilbert space, and $|\hat{f}_k^B\rangle$ is an orthonormal basis of Bob's Hilbert space. The λ_k's are the eigenvalues of Bob's two reduced density matrices corresponding to $|0\rangle$ and $|1\rangle$ (which are identical). There always exists a unitary transformation U that maps an orthonormal basis $|\hat{e}_k^A\rangle$ of a Hilbert space to another orthonormal basis $|\hat{e}_k'^A\rangle$ of the same Hilbert space, and thus this *local* unitary transformation (a rotation on Alice's side only) can map $|0\rangle$ to $|1\rangle$.

Therefore, Alice can start her commit phase with the bit $b = 0$. She prepares the state $|0\rangle$, skips the measurement (delays until Step (4)) and sends Bob's part of the state $|0\rangle$ directly to Bob. At the beginning of the unveil phase, Alice has to choose the value b. If she chooses $b = 0$, she can proceed with the original protocol honestly. If she chooses $b = 1$, she can execute the unitary transformation U, and switch $|0\rangle$ to $|1\rangle$. Bob has no chance to detect the cheating, since his reduced density matrix is the same in both cases.

In the second case, let the two possible reduced density matrices of Bob, corresponding to the two states $|0\rangle$ and $|1\rangle$, be different. They must, however, be similar; otherwise Bob

could easily distinguish between the bits 0 and 1, and so he could cheat. Alice can again use her cheating strategy from above. Mayers [1997] has shown that with a cheating Alice, the probability of Bob being able to distinguish between 0 and 1 will not be larger. Thus, Alice can cheat again with a probability close to 1.

As we can see, a dishonest party can use a *local* action for subsequent modification of the committed bit. Hence, it is impossible for the honest party to detect the cheater, and thus secure quantum bit commitment is not possible.

6. OUTLOOK AND CONCLUSIONS

The security of quantum key distribution relies on the inviolable laws of quantum mechanics: nonorthogonal quantum states are used as signal states in the BB84 protocol. The impossibility of perfect cloning of nonorthogonal states implies the security of this protocol.

In the security proof for the BB84 protocol, we have employed an equivalent entanglement-based protocol. The main idea is that local measurements on a maximally entangled state, shared by Alice and Bob, have perfectly correlated outcomes that can be used as the key. A maximally entangled state is necessarily pure, and a pure state cannot be entangled with an eavesdropper's state—thus Eve cannot learn anything about the key. The idea for quantum cryptography with entangled states goes back to Ekert [1991], who suggested to confirm the existence of quantum correlations in the state of Alice and Bob by a Bell inequality test.

6.1. Other Quantum Key Distribution Protocols

A variety of quantum key distribution protocols can be found in the literature. All known prepare-and-measure schemes can be seen as variations of the BB84 protocol, which are obtained by changing the number and/or dimension of the quantum states.

Bennett [1992] proposed a protocol—which now is named after him the B92 protocol—in which only two nonorthogonal states are used. In the so-called six-state protocol [Bruß 1998; Bechmann-Pasquinucci and Gisin 1999], the six eigenstates of the three Pauli operators are used. In this protocol, it is more difficult for Eve to retrieve any information, thus the security is enhanced.

In this article, we have always considered qubits, that is, two-level systems as information carriers. What happens if one considers higher-dimensional systems, such as qutrits (three-level systems)? Intuitively, one would expect that the increased number of degrees of freedom makes it more difficult for Eve to extract information on the key. As proven in Bruß and Macchiavello [2002], higher-dimensional systems indeed offer increased security.

A recently suggested protocol [Scarani et al. 2004] introduces a new sifting method: rather than announcing the basis, Alice gives Bob a list of two nonorthogonal states from which the signal state was taken. This protocol has certain security advantages that are connected with experimental implementations of quantum cryptography.

6.2. Experimental Status

So far, we have presented quantum key distribution in a rather theoretical, abstract manner. What is the experimental situation—can the ideas of quantum cryptography be made reality? In recent years, much effort has been devoted to experiments on quantum cryptography, and much progress has been made. In most experiments, polarized photons are representing the qubits: photons are polarized if their electromagnetic field oscillates in a fixed direction of space (which has to be orthogonal to the direction of flight). The two degrees of freedom for a photonic qubit can be, for example, horizontal

and vertical polarization (the rectilinear basis in the BB84 protocol), or polarization rotated by $45°$ with respect to the horizontal/vertical direction—this corresponds to the diagonal basis in BB84. The experimentalist "only" has to produce single polarized photons on demand.

This, however, is one of the main experimental challenges: an attenuated laser pulse consists of Poisson-distributed number states, that is, with a certain probability more than one photon will be emitted. These events with more than one photon allow for a dangerous eavesdropping strategy, the so-called photon-number splitting attack, where Eve splits off a photon and receives full information about the key. Apart from experimental progress towards true single-photon sources, new algorithms that can cope with this sort of attack have been developed. One example, the protocol by Scarani et al. [2004], has already been mentioned above. Another important contribution is the so-called decoy state protocol introduced by Hwang [2003], which uses two photon sources with different number statistics to "decoy" the adversary.

The long-term goal in experimental quantum key distribution is to reach high key rates over large distances. For the transmission of photons, two possibilities exist: either transmission via optical fibers, or transmission in free space. Rather than trying to summarize all existing experiments, let us mention just two examples. A very stable, robust system with optical fiber transmission has been developed by Gisin and Zbinden at the University of Geneva [Gisin et al. 2002]. They were able to transmit a secret key from Geneva to Lausanne (i.e., over a distance of about 67 km), with a rate of 130 bit/s. Regarding free space quantum cryptography, Weinfurter from LMU Munich [Kurtsiefer et al. 2002] recently demonstrated secret key exchange over about 23.4 km (in the Alps, from Zugspitze to Karwendelspitze), with a rate of about 1000 bit/s.

For realistic implementations, the above security proof of an ideal protocol does not necessarily hold, due to imperfections in the source (multi-photon signals) and detectors (noise, losses), see Inamori et al. [2001]. Also note that we considered only optimal error correcting codes in our security proof, whereas the error correcting schemes used in practice usually are less efficient. With present technology it is possible to implement unconditionally secure quantum key distribution protocols for distances around 20 km, without using the decoy state method, and for higher distances with decoy pulses.

Long-term goals of quantum key distribution are the realistic implementation via fibers, for example, for different buildings of a bank or company (with a relatively small distance), and free space key exchange via satellites. Future practical developments will have to prove which one of the described protocols will turn out to be successful. At the moment, demonstrators for long-range quantum key distribution are being built within the EU project SECOQC (for further information, see www.secoqc.net). Quantum cryptography already provides the most advanced technology of quantum information science, and is on the way to achieve the (quantum) jump from university laboratories to the real world.

ACKNOWLEDGMENTS

REFERENCES

AJTAI, M. AND DWORK, C. 1997. A public-key cryptosystem with worst-case/average-case equivalence. In *Proceedings of the 29th ACM Symposium on Theory of Computing*. ACM, New York, 284–293.

ALLENDER, E. AND RUBINSTEIN, R. 1988. P-printable sets. *SIAM J. Comput. 17*, 6, 1193–1202.

BECHMANN-PASQUINUCCI, H., AND GISIN, N. 1999. Incoherent and coherent eavesdropping in the six-state protocol of quantum cryptography. *Phys. Rev. A 59*, 4238–4248.

BEN-OR, M., HORODECKI, M., LEUNG, D., MAYERS, D., AND OPPENHEIM, J. 2005. The universal composable security of quantum key distribution. In *Proceedings of the 2nd Theory of Cryptography Conference*, J. Kilian, Ed. Lecture Notes in Computer Science, vol. 3378, Springer-Verlag, 386–406. Also available at http://arxiv.org/abs/quant-ph/0409078.

BENNETT, C. 1992. Quantum cryptography using any two nonorthogonal states. *Phys. Rev. Lett. 68*, 3121–3124.

BENNETT, C. AND BRASSARD, G. 1984. Quantum cryptography: Public key distribution and coin tossing. In *Proceedings of the IEEE International Conference on Computers, Systems, and Signal Processing*. IEEE Computer Society Press, Los Alamitos, CA, 175–179.

BENNETT, C., BRASSARD, G., CRÉPEAU, C., AND MAURER, U. 1995. Generalized privacy amplification. *IEEE Trans. Inf. Theory 41*, 1915–1923.

BENNETT, C., BRASSARD, G., AND MERMIN, D. 1992. Quantum cryptography without Bell's theorem. *Phys. Rev. Lett. 68*, 557–559.

BERMAN, L. 1977. Polynomial reducibilities and complete sets. Ph.D. dissertation, Cornell University, Ithaca, NY.

BEYGELZIMER, A., HEMASPAANDRA, L., HOMAN, C., AND ROTHE, J. 1999. One-way functions in worst-case cryptography: Algebraic and security properties are on the house. *SIGACT News 30*, 4 (Dec.), 25–40.

BOUWMEESTER, D., EKERT, A., AND ZEILINGER, A. 2000. *The Physics of Quantum Information*. Springer-Verlag, New York.

BRASSARD, G. 1979. A note on the complexity of cryptography. *IEEE Trans. Inf. Theory 25*, 2, 232–233.

BRASSARD, G., CRÉPEAU, C., JOZSA, R., AND LANGLOIS, D. 1993. A quantum bit commitment scheme provably unbreakable by both parties. In *Proceedings of the 34th IEEE Symposium on Foundations of Computer Science*. IEEE Computer Society Press, Los Alamitos, CA, 362–371.

BRASSARD, G., FORTUNE, S., AND HOPCROFT, J. 1978. A note on cryptography and NP ∩ coNP − P. Tech. Rep. TR-338, Department of Computer Science, Cornell University, Ithaca, NY. Apr.

BRUSS, D. 1998. Optimal eavesdropping in quantum cryptography with six states. *Phys. Rev. Lett. 81*, 3018–3021.

BRUSS, D., AND MACCHIAVELLO, C. 2002. Optimal eavesdropping in cryptography with three-dimensional quantum states. *Phys. Rev. Lett. 88*, 127901(1)–127901(4).

DAMGÅRD, I., FEHR, S., SALVAIL, L., AND SCHAFFNER, C. 2005. Cryptography in the bounded quantum-storage model. In *Proceedings of the 46th IEEE Symposium on Foundations of Computer Science*. IEEE Computer Society Press, Los Alamitos, CA, 449–458.

DIFFIE, W. AND HELLMAN, M. 1976. New directions in cryptography. *IEEE Trans. Inf. Theory IT-22*, 6, 644–654.

EKERT, A. 1991. Quantum cryptography based on Bell's theorem. *Phys. Rev. Lett. 67*, 661–663.

ELGAMAL, T. 1985. A public key cryptosystem and a signature scheme based on discrete logarithms. *IEEE Trans. Inf. Theory IT-31*, 4, 469–472.

FELLOWS, M. AND KOBLITZ, N. 1992. Self-witnessing polynomial-time complexity and prime factorization. *Des. Codes Crypt. 2*, 3, 231–235.

FENNER, S., FORTNOW, L., NAIK, A., AND ROGERS, J. 2003. Inverting onto functions. *Inf. Comput. 186*, 1, 90–103.

GISIN, N., RIBORDY, G., TITTEL, W., AND ZBINDEN, H. 2002. Quantum cryptography. *Rev. Modern Phys. 74*, 145–195.

GOLDREICH, O. 1997. Note on Levin's theory of average-case complexity. Tech. Rep. TR97-058, Electronic Colloquium on Computational Complexity. Nov.

GROLLMANN, J. AND SELMAN, A. 1988. Complexity measures for public-key cryptosystems. *SIAM J. Comput. 17*, 2, 309–335.

HARTMANIS, J. AND HEMACHANDRA, L. 1991. One-way functions and the nonisomorphism of NP-complete sets. *Theoret. Comput. Sci. 81*, 1, 155–163.

HÅSTAD, J., IMPAGLIAZZO, R., LEVIN, L., AND LUBY, M. 1999. A pseudorandom generator from any one-way function. *SIAM J. Comput. 28*, 4 (Aug.), 1364–1396.

HEMASPAANDRA, L., PASANEN, K., AND ROTHE, J. 2006. If P ≠ NP then some strongly noninvertible functions are invertible. *Theoret. Comput. Sci. 362*, 1–3, 54–62.

HEMASPAANDRA, L. AND ROTHE, J. 1999. Creating strong, total, commutative, associative one-way functions from any one-way function in complexity theory. *J. Comput. Syst. Sci. 58*, 3, 648–659.

HEMASPAANDRA, L. AND ROTHE, J. 2000. Characterizing the existence of one-way permutations. *Theoret. Comput. Sci. 244*, 1–2 (Aug.), 257–261.

HEMASPAANDRA, L., ROTHE, J., AND SAXENA, A. 2005. Enforcing and defying associativity, commutativity, totality, and strong noninvertibility for one-way functions in complexity theory. In *Proceedings of the 9th Italian Conference on Theoretical Computer Science*. Lecture Notes in Computer Science, vol. 3701. Springer-Verlag, New York, 265–279.

HEMASPAANDRA, L., ROTHE, J., AND WECHSUNG, G. 1997a. Easy sets and hard certificate schemes. *Acta Inf. 34*, 11 (Nov.), 859–879.

HEMASPAANDRA, L., ROTHE, J., AND WECHSUNG, G. 1997b. On sets with easy certificates and the existence of one-way permutations. In *Proceedings of the 3rd Italian Conference on Algorithms and Complexity*. Lecture Notes in Computer Science, vol. 1203. Springer-Verlag, New York, 264–275.

HOMAN, C. 2004. Tight lower bounds on the ambiguity of strong, total, associative, one-way functions. *J. Comput. Syst. Sci. 68*, 3, 657–674.

HOMAN, C. AND THAKUR, M. 2003. One-way permutations and self-witnessing languages. *J. Comput. Syst. Sci. 67*, 3, 608–622.

HUFFMAN, W., AND PLESS, V. 2003. *Fundamentals of Error-Correcting Codes*. Cambridge University Press, Cambridge, MA.

HWANG, W. 2003. Quantum key distribution with high loss: Toward global secure communication. *Phys. Rev. Lett. 91*, 057901.

INAMORI, H., LÜTKENHAUS, N., AND MAYERS, D. 2001. Unconditional security of practical quantum key distribution. Tech. Rep. quant-ph/0107017, Computing Research Repository (CoRR). Available on-line at http://arxiv.org/abs/quant-ph/0107017.

KAWACHI, A., KOBAYASHI, H., KOSHIBA, T., AND PUTRA, R. 2005. Universal test for quantum one-way permutations. *Theoret. Comput. Sci. 345*, 370–385.

KENT, A. 1999. Unconditionally secure bit commitment. *Phys. Rev. Lett. 83*, 1447–1450.

KO, K. 1985. On some natural complete operators. *Theoret. Comput. Sci. 37*, 1, 1–30.

KÖNIG, R., RENNER, R., BARISKA, A., AND MAURER, U. 2006. Locking of accessible information and implications for the security of quantum cryptography. Tech. Rep. quant-ph/0512021v2, Computing Research Repository (CoRR). Available on-line at http://arxiv.org/abs/quant-ph/0512021.

KURTSIEFER, C., ZARDA, P., HALDER, M., WEINFURTER, H., GORMAN, P., TAPSTER, P., AND RARITY, J. 2002. A step towards global key distribution. *Nature 419*, 450.

LEVIN, L. 1986. Average case complete problems. *SIAM J. Comput. 15*, 1, 285–286.

LO, H. AND CHAU, H. 1997. Is quantum bit commitment really possible? *Phys. Rev. Lett. 78*, 3410–3413.

LO, H. AND CHAU, H. 1999. Unconditional security of quantum key distribution over arbitrarily long distances. *Science 283*, 2050–2056.

MAURER, U. 1993. Secret key agreement by public discussion from common information. *IEEE Trans. Inf. Theory 39*, 733–742.

MAURER, U. AND WOLF, S. 1999. The relationship between breaking the Diffie-Hellman protocol and computing discrete logarithms. *SIAM J. Comput. 28*, 5, 1689–1721.

MAY, A. 2004. Computing the RSA secret key is deterministic polynomial time equivalent to factoring. In *Advances in Cryptology—CRYPTO '04*. Lecture Notes in Computer Science, vol. 3152. Springer-Verlag, New York, 213–219.

MAYERS, D. 1995. The trouble with quantum bit commitment. Tech. Rep. quant-ph/9603015v3, Computing Research Repository (CoRR). Available on-line at http://arxiv.org/abs/quant-ph/9603015.

MAYERS, D. 1997. Unconditionally secure quantum bit commitment is impossible. *Phys. Rev. Lett. 78*, 3414–3417.

NGUYEN, P. AND STERN, J. 1998. Cryptanalysis of the Ajtai-Dwork cryptosystem. In *Advances in Cryptology—CRYPTO '98*. Lecture Notes in Computer Science, vol. 1462. Springer-Verlag, New York, 223–242.

NIELSEN, M. AND CHUANG, I. 2000. *Quantum Computation and Quantum Information*. Cambridge University Press, Cambridge, MA.

RABI, M. AND SHERMAN, A. 1997. An observation on associative one-way functions in complexity theory. *Inf. Proc. Lett. 64*, 5, 239–244.

RIVEST, R., SHAMIR, A., AND ADLEMAN, L. 1978. A method for obtaining digital signature and public-key cryptosystems. *Commun. ACM 21*, 2, 120–126.

ROTHE, J. 2005. *Complexity Theory and Cryptology. An Introduction to Cryptocomplexity*. EATCS Texts in Theoret. Comput. Sci. Springer-Verlag, Berlin, Heidelberg, New York.

ROTHE, J. AND HEMASPAANDRA, L. 2002. On characterizing the existence of partial one-way permutations. *Inf. Proc. Lett. 82*, 3 (May), 165–171.

SCARANI, V., ACÍN, A., RIBORDY, G., AND GISIN, N. 2004. Quantum cryptography protocols robust against photon number splitting attacks for weak laser pulse implementations. *Phys. Rev. Lett. 92*, 057901(1)–057901(4).

SCHRÖDINGER, E. 1935. Die gegenwärtige Situation in der Quantenmechanik. *Die Naturwissenschaften 23*, 807–812, 823–828, 844–849.

SELMAN, A. 1992. A survey of one-way functions in complexity theory. *Math. Syst. Theory 25*, 3, 203–221.

SHANNON, C. 1949. Communication theory of secrecy systems. *Bell Syst. Tech. J. 28*, 4, 657–715.

SHOR, P. 1997. Polynomial-time algorithms for prime factorization and discrete logarithms on a quantum computer. *SIAM J. Comput. 26*, 5, 1484–1509.

SHOR, P. AND PRESKILL, J. 2000. Simple proof of security of the BB84 quantum key distribution protocol. *Phys. Rev. Lett. 85*, 441–444.

SINGH, S. 1999. *The Code Book. The Science of Secrecy from Ancient Egypt to Quantum Cryptography*. Fourth Estate, London, England.

STINSON, D. 2005. *Cryptography: Theory and Practice*, Third ed. CRC Press, Boca Raton, FL.

WANG, J. 1997. Average-case computational complexity theory. In *Complexity Theory Retrospective II*, L. Hemaspaandra and A. Selman, Eds. Springer-Verlag, New York, 295–328.

WATANABE, O. 1988. On hardness of one-way functions. *Inf. Proc. Lett. 27*, 3, 151–157.

WERNER, R. 1989. Quantum states with Einstein–Podolsky–Rosen correlations admitting a hidden-variable model. *Phys. Rev. A 40*, 8 (Oct.), 4277–4281.

WIESNER, S. 1983. Conjugate coding. *SIGACT News 15*, 78–88.

WOOTTERS, W. K., AND ZUREK, W. H. 1982. A single quantum cannot be cloned. *Nature 299*, 802–803.

YAO, A. 1995. Security of quantum protocols against coherent measurements. In *Proceedings of the 27th ACM Symposium on Theory of Computing*. ACM, New York, 67–75.